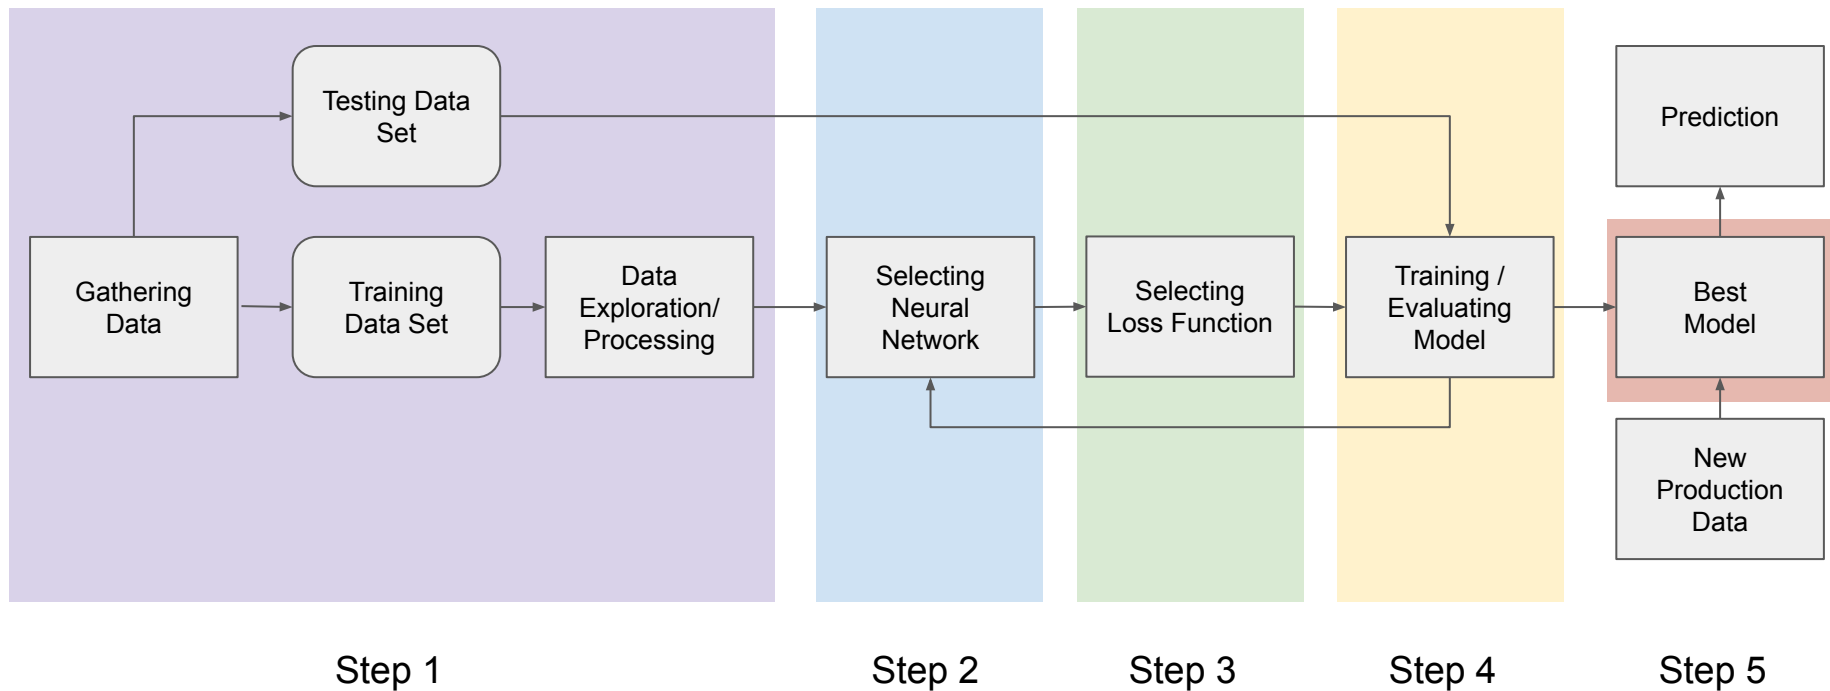


# Learning Deep Learning with PyTorch

## (2) Mechanics of Learning

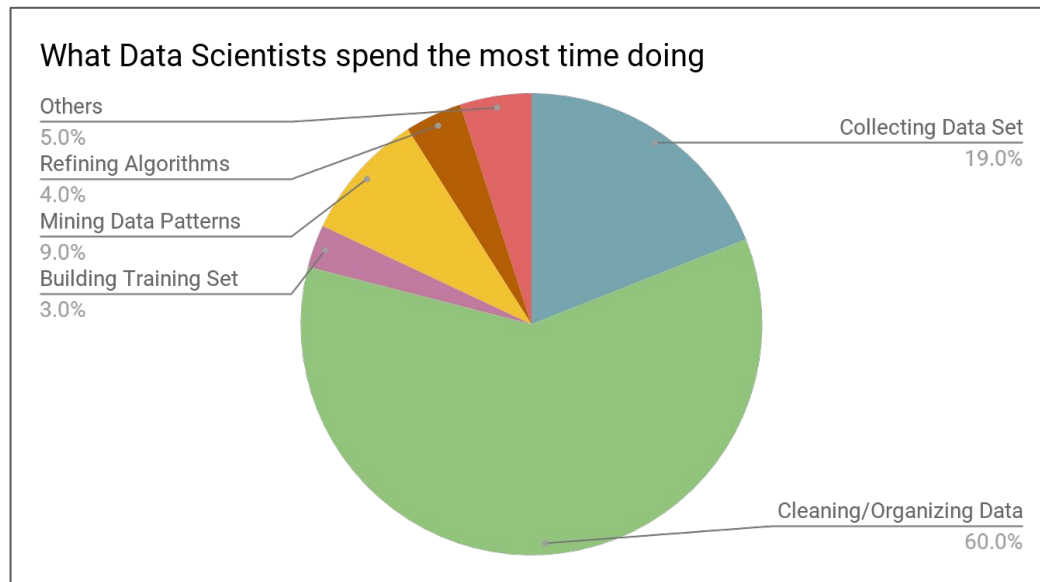
Qiyang Hu  
UCLA IDRE  
April 20, 2020

# Workflow for a deep learning project



# Step 1. Data Prep

- The most time-consuming but the most *creative* job
  - Take > 80% time
  - Require experience
  - May need domain expertise
- Determines the upper limit for the goodness of DL
  - Models/Algorithms: just approach the upper limit

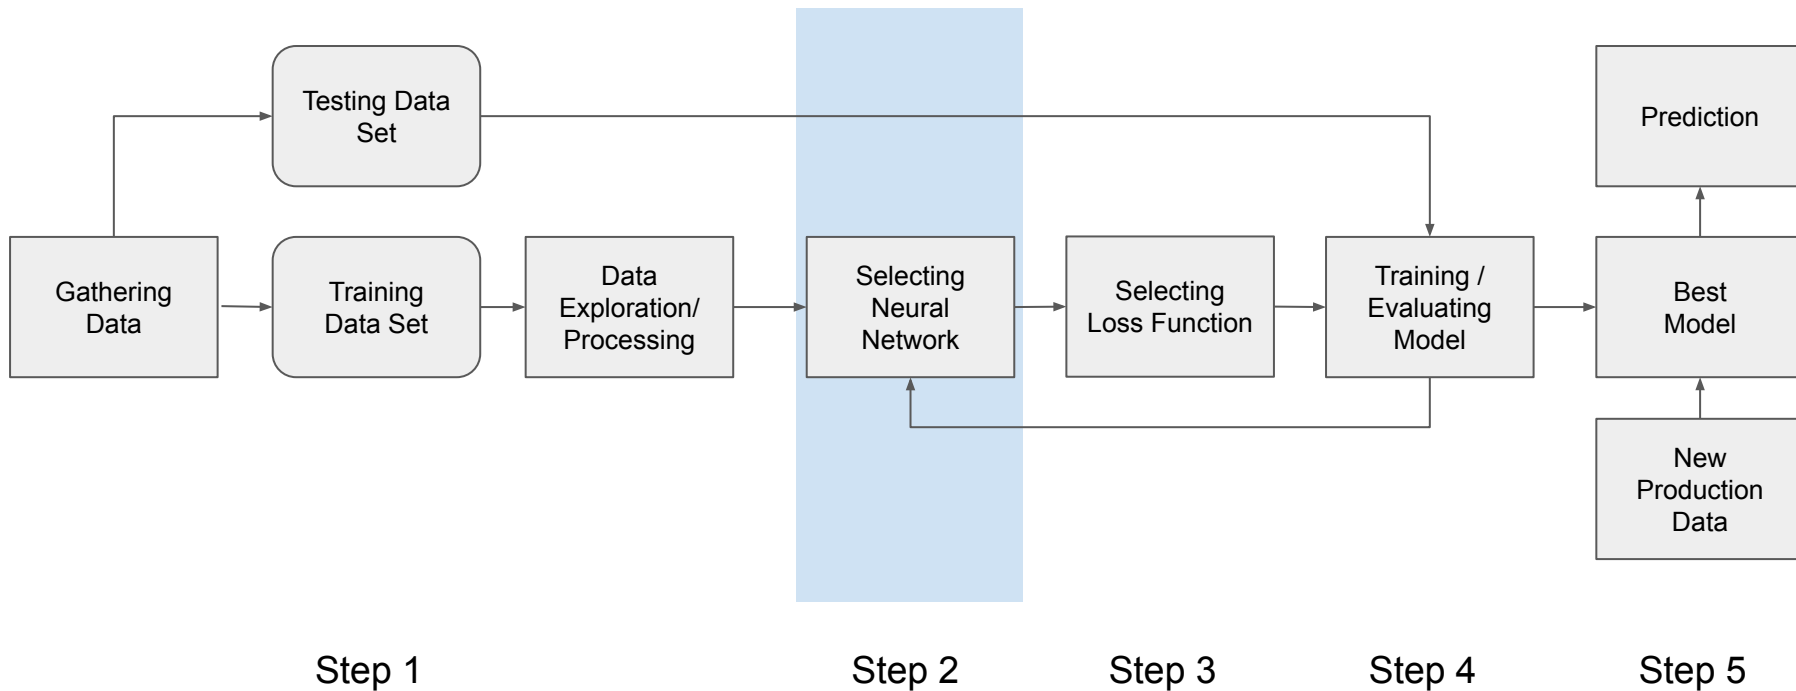


Survey from Forbes in 2017 ([Data Source](#))

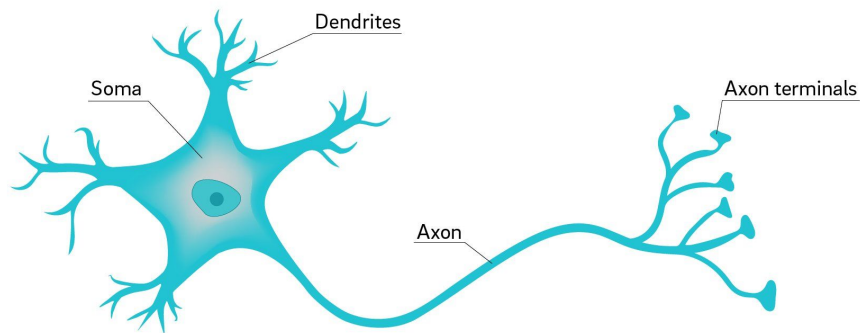
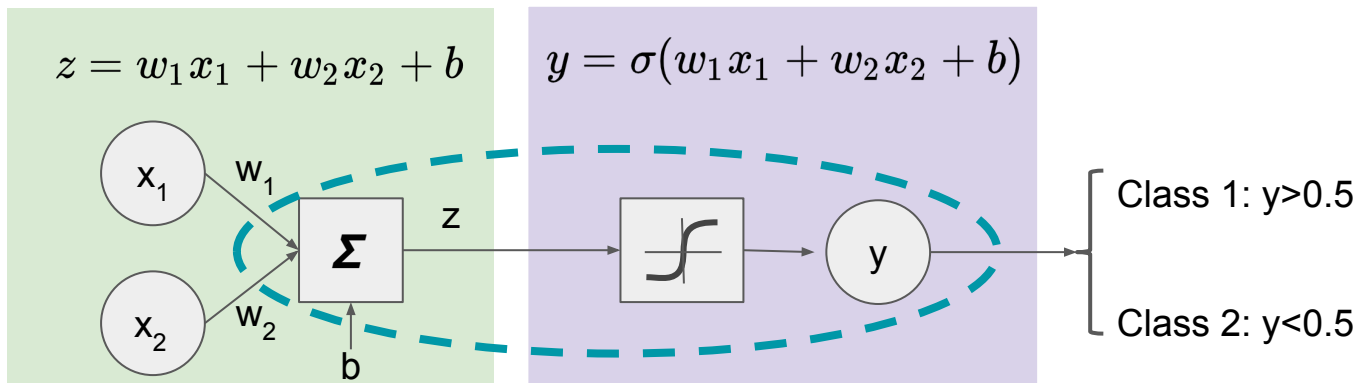
# Feature Engineering

- Transforming raw data into features with a good representation
  - Deep learning can extract hierarchical features automatically
  - DL still needs the digitalization of the input raw data and some form of prior knowledge.
- Some common techniques:
  - Imputation (almost every column)
  - Label binarization (e.g. sex)
  - One-hot encoding (nominal categorical data)
  - Binning and grouping
  - Scaling: standardization and normalization (numerical values with different ranges)
  - Splitting the features (e.g. title from name, etc)


# Workflow for a deep learning project



# Recap: A linear classifier ~ one artificial neuron

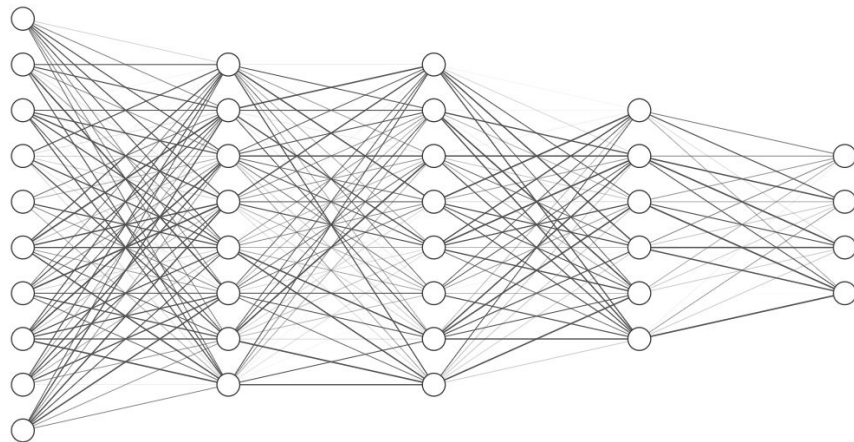



# (Deep) Neural Networks ~ piling/stacking logistic-regression classifiers



# Why deep?

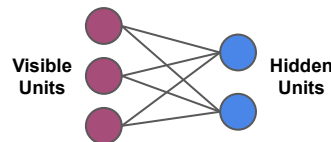

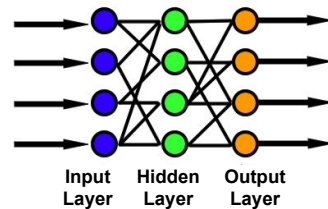
- Shallow network can fit any function
  - Has less number of hidden layers
  - Has to be really “fat”
- Deep network is more efficient.
  - It can extract/build better features
  - Exponentially fewer parameters ([2017](#))





# Types of Neural Network Architectures

- **Feed forward neural networks** (No cycle in node connections)
  - Fully connected network
  - Convolutional networks (CNNs)
- **Recurrent networks** (w/ directed cycle in node connections)
  - Fully recurrent NN
  - Recursive NN
  - Long short-term memory (LSTM)
  - Hopfield network (w/o hidden nodes)
- **Symmetric networks** (no directions in node connections)
  - Boltzmann Machines
    - RBM, DBM



# Activation Function

- Sigmoid function:  $\sigma(z) = \frac{1}{1 + \exp(-z)}$
- tanh function:  $\tanh(z) = \frac{\exp(z) - \exp(-z)}{\exp(z) + \exp(-z)}$

- Rectified linear unit (ReLU)

- Softplus
- Leaky ReLU
- Exponential LU (ELUs)
- GELU, etc.



$$f(x) = x^+ = \max(0, x)$$

- Softmax function:

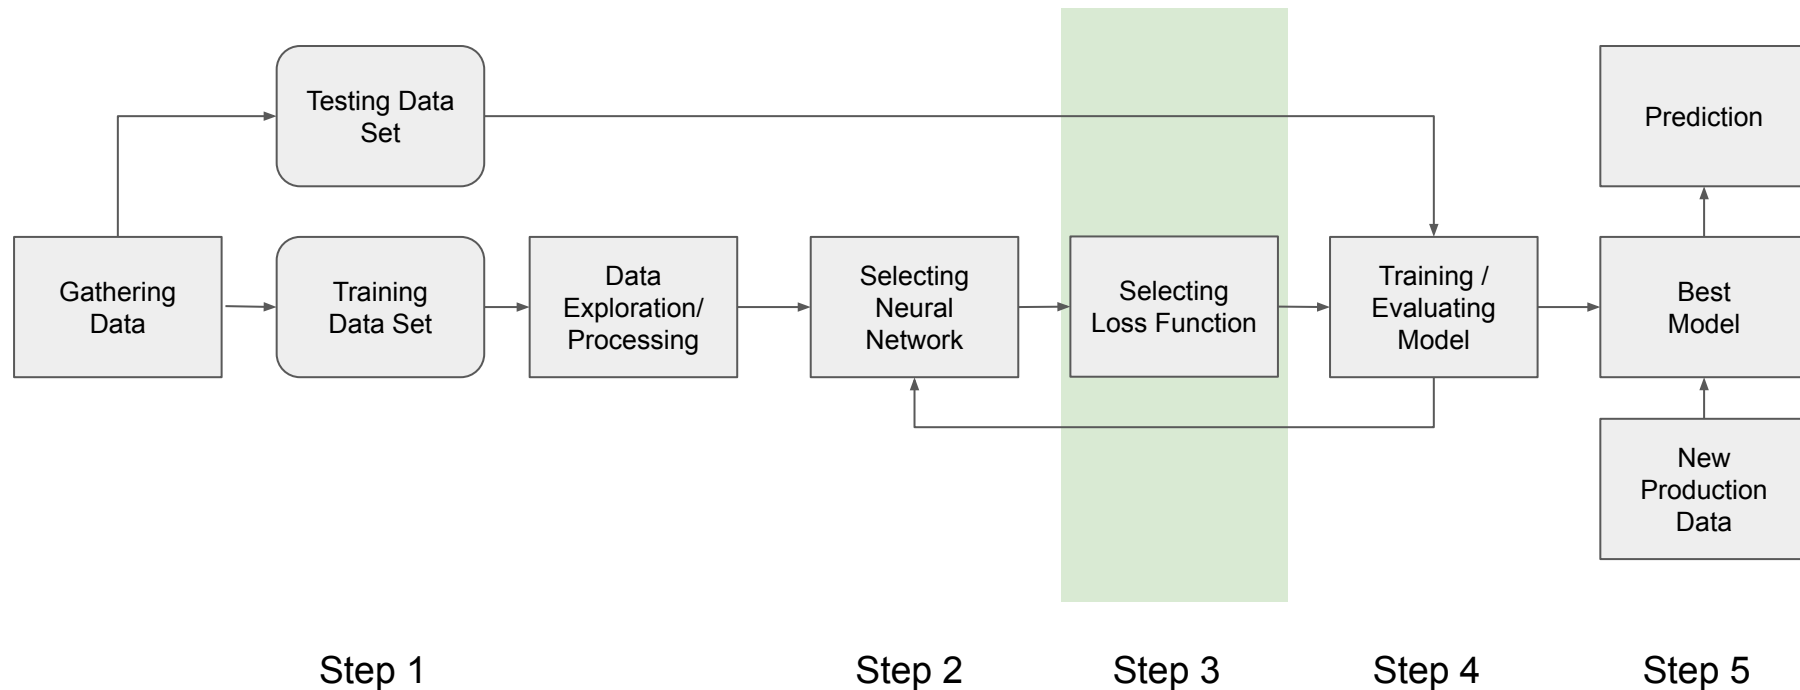
$$y_i = \frac{e^{z^{(i)}}}{\sum_{j=0}^K e^{z^{(j)}}}$$

- Maxout Network:

- *Learnable* activation function



# Workflow for a deep learning project



# How to measure the performance of the model?

- General name: objective function
- Measure the misfit of the model as a function of parameters
  - Criterion is to *minimize* the error functions
  - Loss Function: for a single training example
  - Cost Function: over the entire or mini-batch training set
- Evaluate the probability of generating training set
  - Criterion is to *maximize* the distribution likelihood as a function of parameters
  - Maximum (log)-likelihood estimation
- Regression losses and classification losses

# Loss functions

- Generative/Predictive:



- Regression Loss

- Mean Square Error / Quadratic Loss / L2 Loss:

$$L_{MSE} = \frac{1}{n} \sum_i^n (t_i - s_i)^2$$

- Mean Absolute Error / L1 Loss:

$$L_{MAE} = \frac{1}{n} \sum_i^n |t_i - s_i|$$

- Cross-Entropy Loss and variations


- Softmax Loss / Log Loss / Negative Log Likelihood

- Weighted CE / Balanced CE / Focal Loss

- Dice Loss / IOU Loss / Tversky Loss

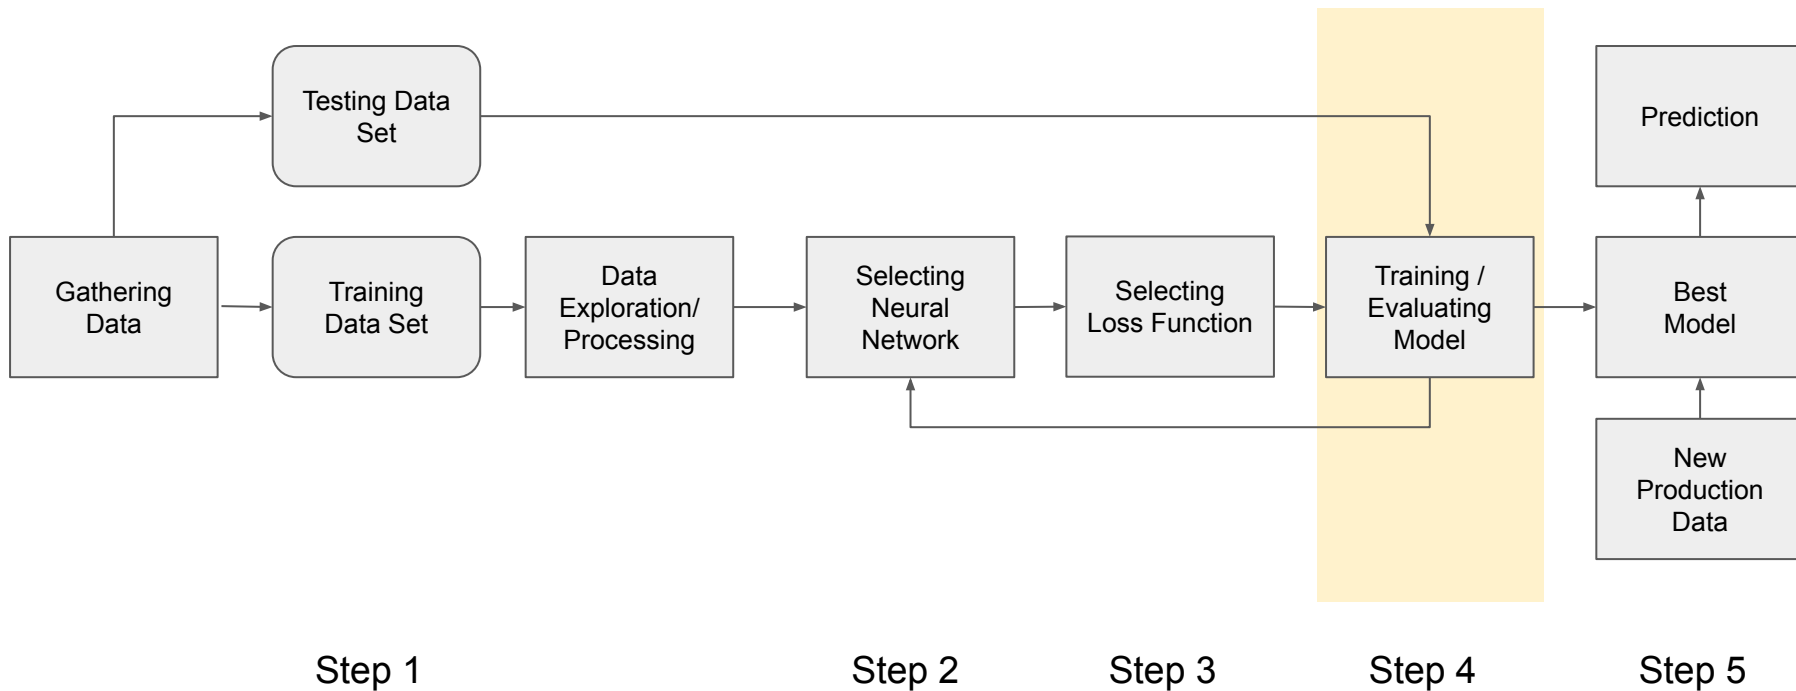
$$L_{CE} = - \sum_i^C t_i \log(s_i)$$

- Contrastive:

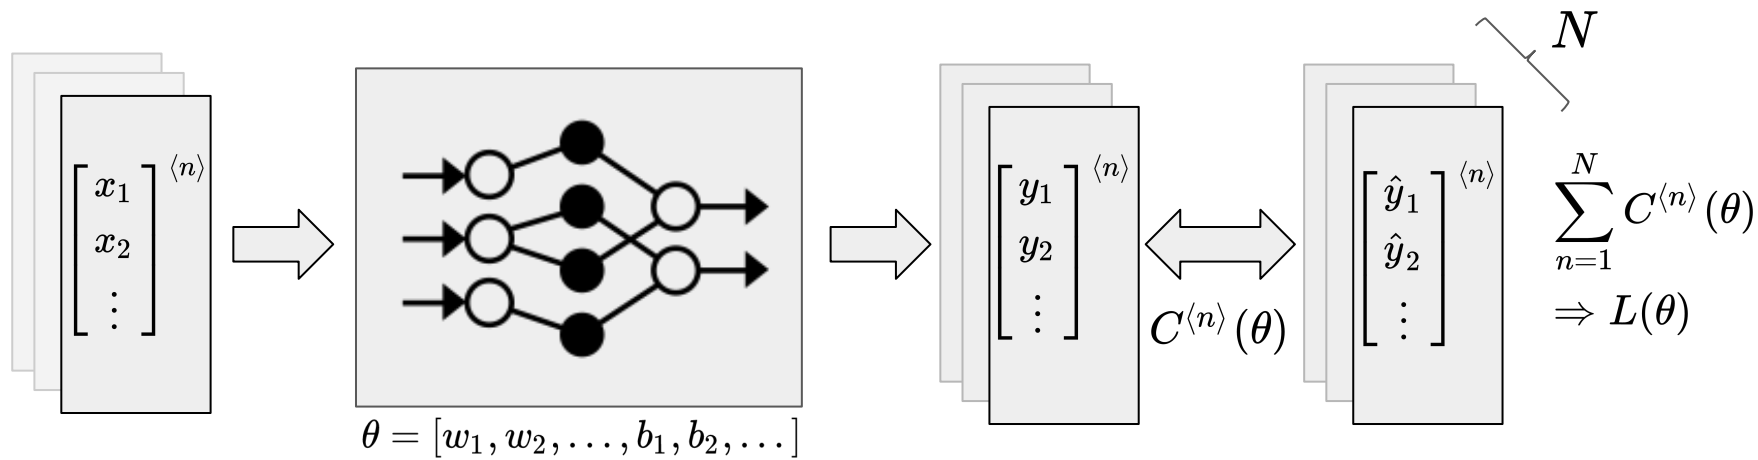


- Ranking Loss/Margin Loss/Contrastive Loss/Triplet Loss

# Workflow for a deep learning project



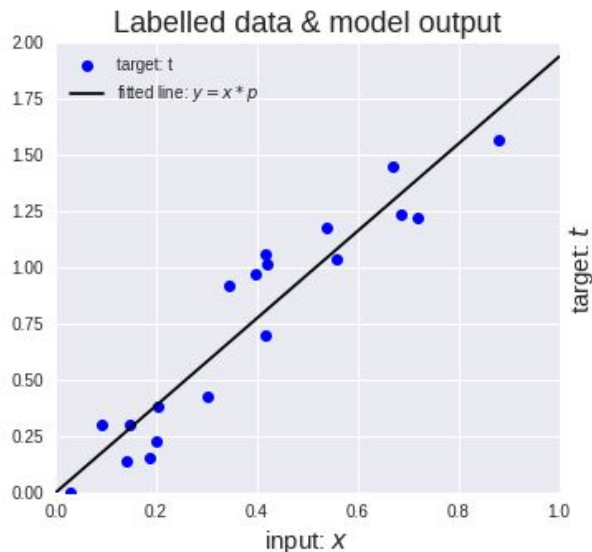

# Training a DNN is an optimization problem



- We know how to compute  $L(\theta)$ , analytically or numerically.
- Start from an arbitrary initialization of  $\theta_o$ , and get an initial  $L_o(\theta)$
- Apply optimization algorithm to minimize  $L(\theta)$

# DL Optimization Algorithm

- Gradient Descent (a 1st-order approach)  $\theta \leftarrow \theta - \eta \nabla L(\theta)$ 
  - Most popular algorithm
    - Pros: simple and fast
    - Cons: sometimes hard to tune

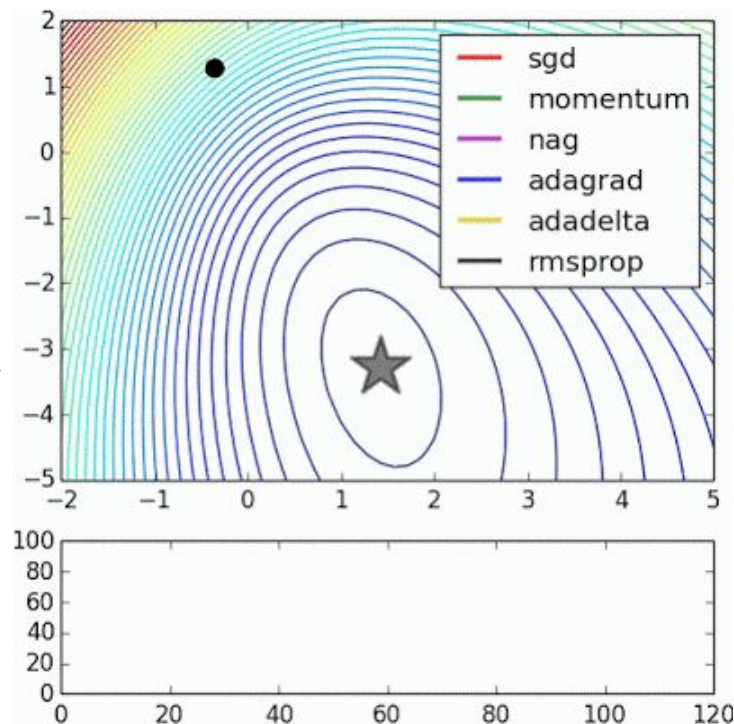


[Source Link](#)



# Gradient-Descent Optimizers

- Stochastic GD / Mini-Batch GD
- Adding momentum:
  - Classical Momentum (CM)
  - Nesterov's Accelerated Gradient (NAG)
- Adaptive learning rate:
  - AdaGrad, AdaDelta, ...
  - RMSprop (animation [source](#)) →
- Combining the two
  - **ADAM** (as **default** in many libs)
- Beyond Adam:
  - Lookahead ([2019](#))
  - RAdam ([2019](#))
  - AdaBound/AmsBound ([ICLR 2019](#))
  - Range ([2019](#))



# Higher Order Optimization Algorithm

- Newton-like methods (2nd-order methods)

$$\theta \leftarrow \theta - \frac{\ell'(\theta)}{\ell''(\theta)}$$

- Prons:
  - Fewer iterations (quadratic convergence)
  - Fewer hyperparameters
- Cons:
  - Much more **costly** in each iteration
  - Need more storing
- BFGS/L-BFGS: a quasi-newton one
  - Good for low dimensional models
- CG (Conjugate gradient)
  - moderately high dimensional models

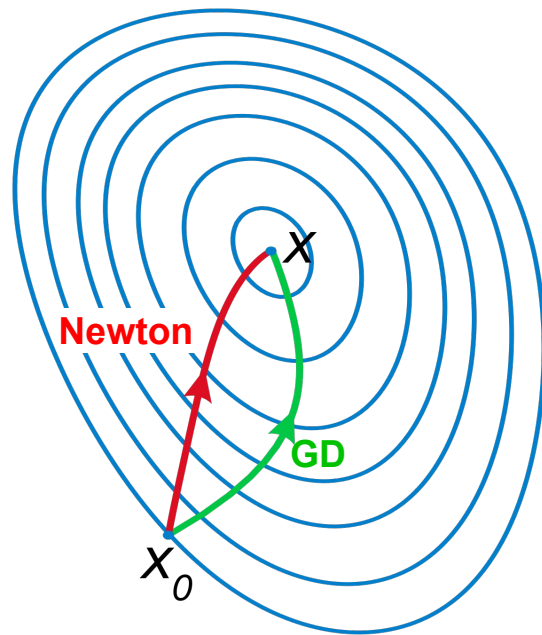
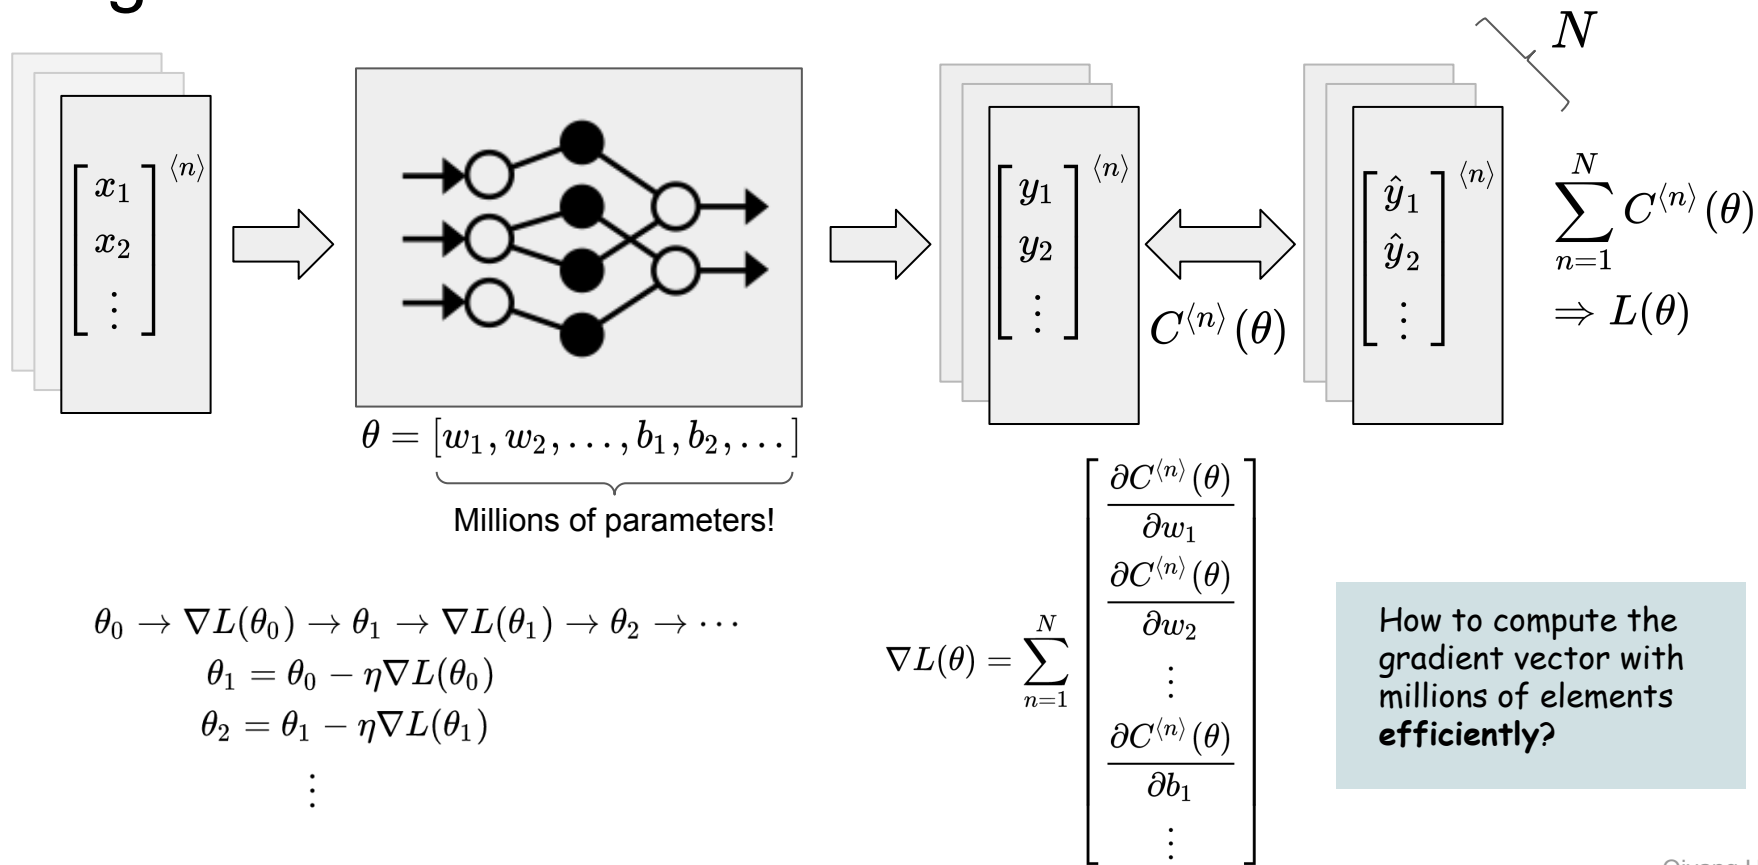



Figure from [Wikipedia](https://en.wikipedia.org/wiki/Newton%27s_method)

# Using Gradient Descent to train DNN



# Backpropagation: a game of chain rule



$$y = \sigma_L \left( w_L \cdot \sigma_{L-1} \left( \cdots w_2 \cdot \sigma_1 \left( \underbrace{w_1 \cdot x + b_1}_{z_1} \right) + b_2 \right) + b_L \right)$$

$$\frac{\partial C(y(w) - \hat{y})}{\partial w} = \frac{\partial z}{\partial w} \frac{\partial C}{\partial z} = \frac{\partial z}{\partial w} \left[ \frac{\partial a}{\partial z} \frac{\partial C}{\partial a} \right] = \frac{\partial z}{\partial w} \left[ \sigma' \cdot \left( \underbrace{\frac{\partial z_{(+1)}}{\partial a} \frac{\partial C}{\partial z_{(+1)}}}_{a_1} \right) \right]$$

## ① Forward Pass


$$\frac{\partial z_1}{\partial w_1} = x_1 \longrightarrow \frac{\partial z_2}{\partial w_2} = a_1 \longrightarrow \cdots \longrightarrow \frac{\partial z_{L-1}}{\partial w_{L-1}} = a_{L-2} \longrightarrow \frac{\partial z_L}{\partial w_L} = a_{L-1}$$

## ② Backward Pass

$$\frac{\partial C}{\partial z_1} = \sigma'_1 \left[ w_2 \frac{\partial C}{\partial z_2} \right] \longleftarrow \cdots \longleftarrow \frac{\partial C}{\partial z_{L-1}} = \sigma'_{L-1} \left[ w_L \frac{\partial C}{\partial z_L} \right] \longleftarrow \frac{\partial C}{\partial z_L} = \sigma'_L \frac{\partial C}{\partial y} \longleftarrow \frac{\partial C}{\partial y}$$

# Wait, here is a catch...

- ReLU as one of the most popular activation functions:  $f(x) = x^+ = \max(0, x)$
- ReLU is not differentiable at  $x=0$
- Why we can use it in gradient based DNN training?
  - NN training *rarely* arrives at a local minimum of the cost function
  - Software implementations of NN training usually return one of the one-sided derivatives (*sub-gradient*)
- In practice we can safely disregard the non-differentiability of the hidden unit activation functions.



# Workflow for a deep learning project

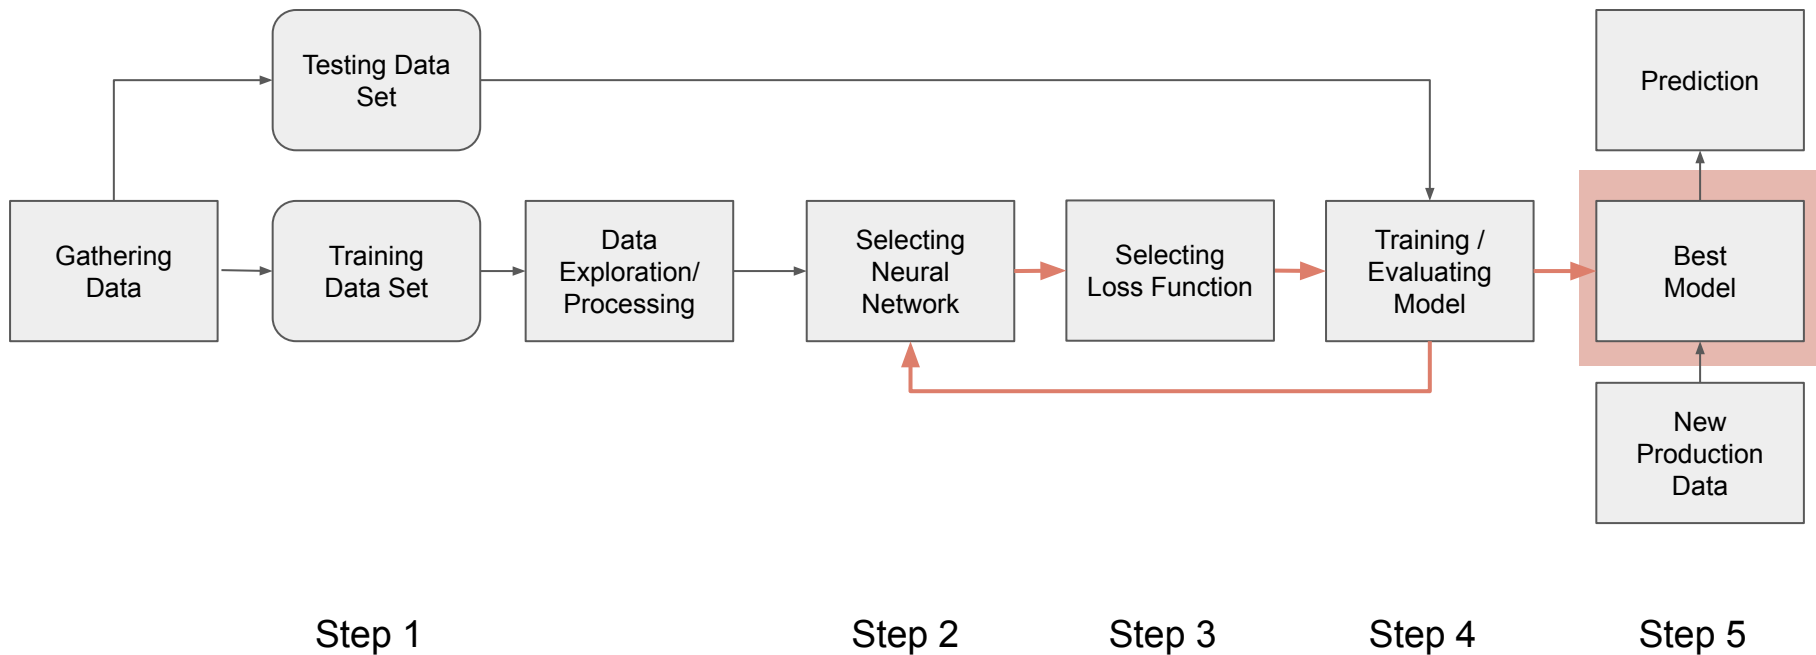
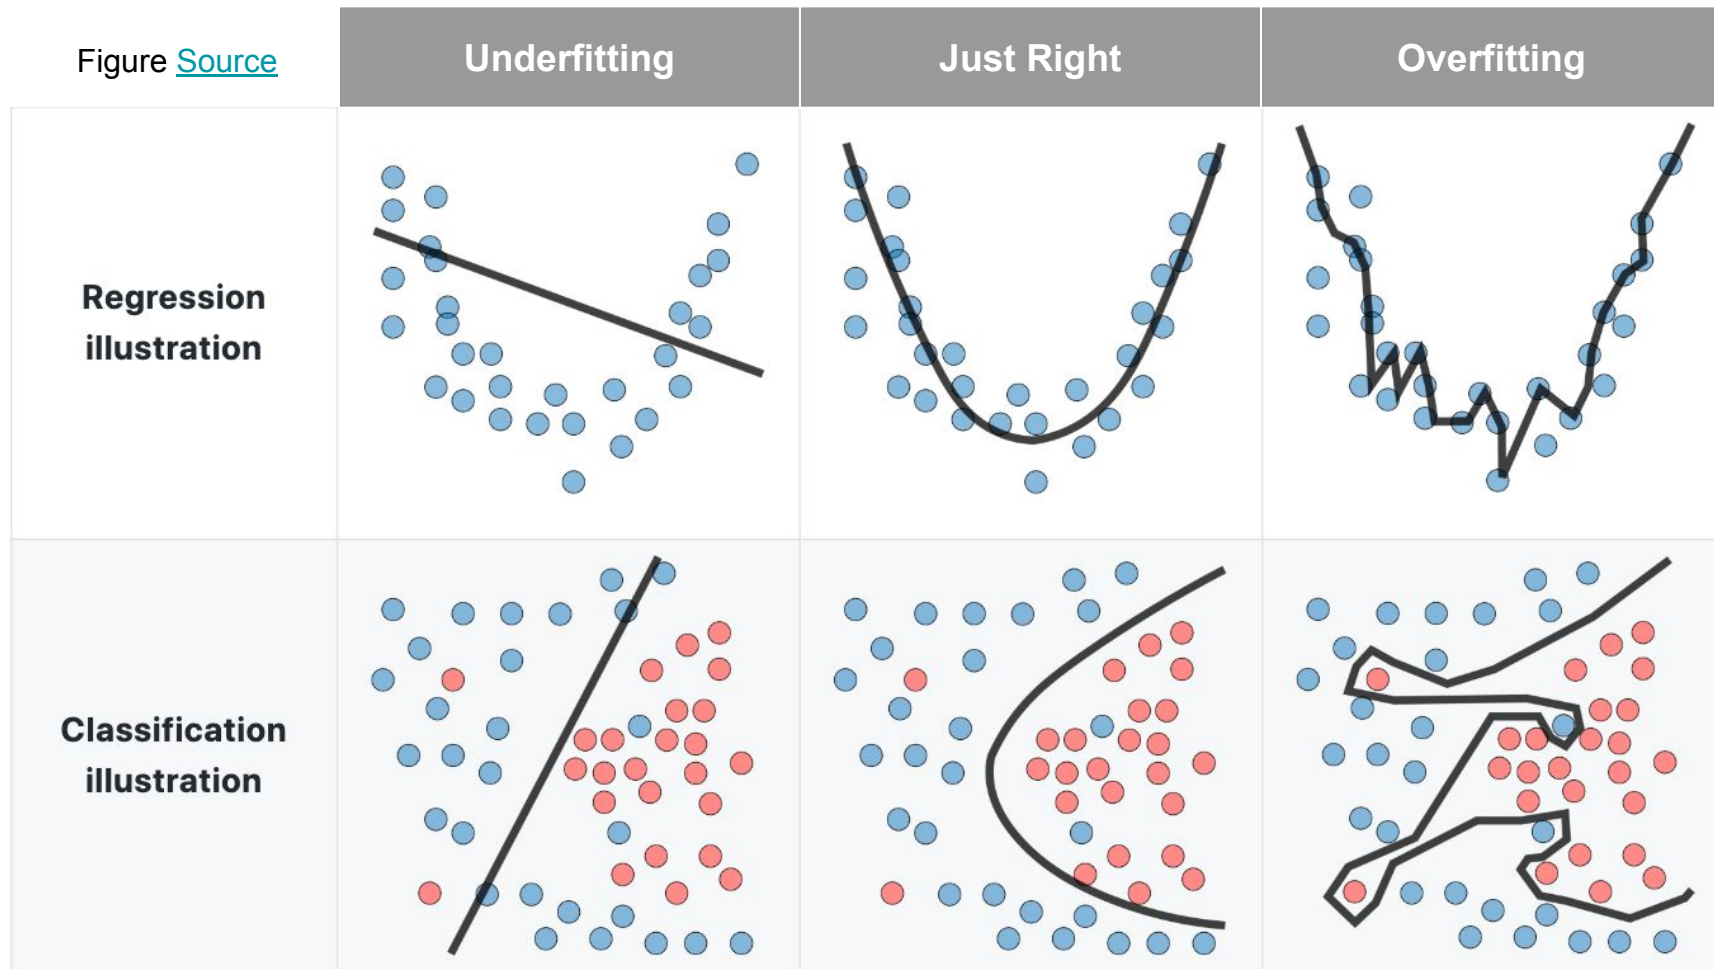



Figure [Source](#)



# Underfitting and Overfitting

- Underfitting: model too simple:
  - Diagnose:
    - cannot even fit the training data
    - training error ~ testing error
  - Ignore the variance in training data
  - Higher prediction bias
- Overfitting: model too complex
  - Diagnose:
    - well-fit for training data
    - large error for testing data
  - Over-interpret training data
  - More deviation from new data





# How to prevent underfitting?

- Redesign the model
- Increase model's complexity
- Add more features as input
- Training longer
- More data will not help

# How to prevent overfitting?

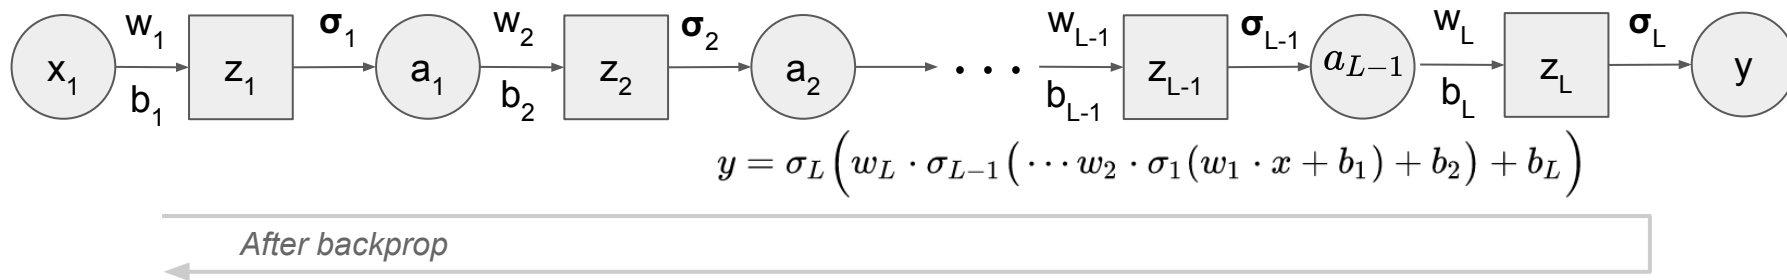
- Get more data
  - Collect more data
  - Data augmentation
- Reduce the model's complexity
- Regularization
  - Weight Regularization to make the model smoother (L1, L2, Elastic net)

$$\hat{L}(x, y) = L(x, y) + \lambda \sum_{i=1}^n \theta_i^2$$

- Early stopping

# Gradient vanishing/exploding in DL training

- Causes

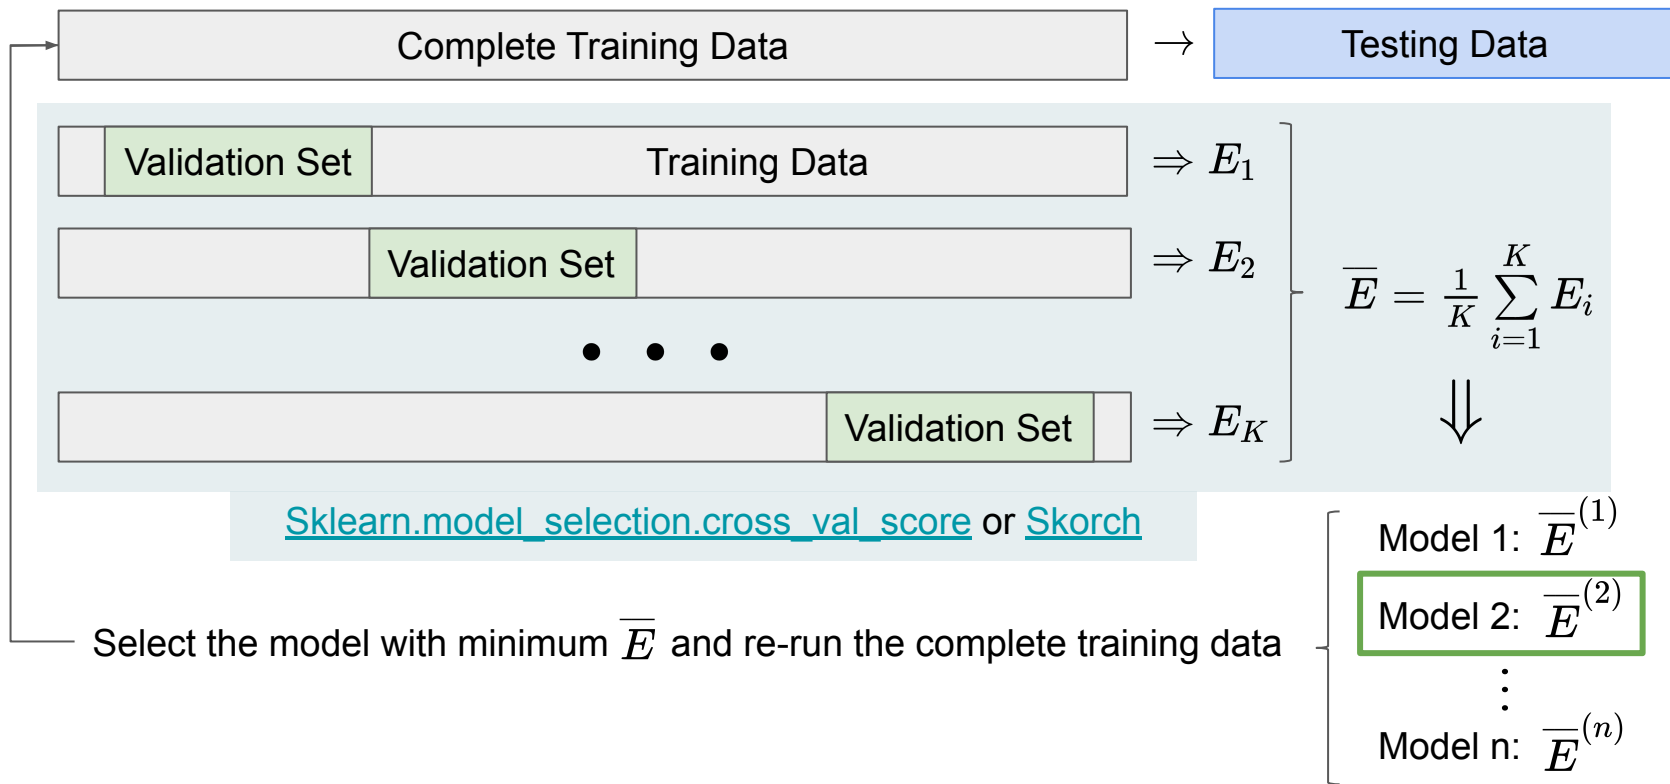


- Gradients in initial layers = Multiplication of Gradients at prior layers
- Small variation around 1 results in vanishing/exploding

- Techniques to resolve:

- General: adjusting learning rate, dropout, batch normalization, layer normalization
- For gradient exploding: gradient clipping, weight regularization
- For gradient vanishing: activation function, proper initialization parameters, LSTM, skip connections

# Model Selection: K-fold Cross Validation



# Errors/scores in practice



Error:  $E^{val} < E^{Pub} < E^{Pri}$

Score:  $S^{val} > S^{Pub} > S^{Pri}$

# Don't forget to

- Github Repo:
  - <https://github.com/huqy/idre-learning-deep-learning-pytorch>
- Slack workspace:
  - [bit.ly/join-LDL](https://bit.ly/join-LDL)
- Contact me
  - [huqy@idre.ucla.edu](mailto:huqy@idre.ucla.edu)
  - Direct message in Slack
- IF you don't have plan to attend the rest of workshops, :
  - Please fill out our series survey: [bit.ly/2X2phyS](https://bit.ly/2X2phyS)