

Learning Deep Learning with PyTorch

(6) RNNs and LSTM

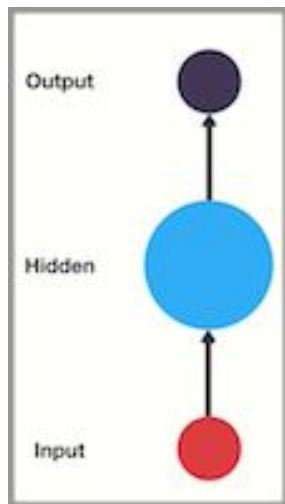
Qiyang Hu

IDRE

May 7th, 2020

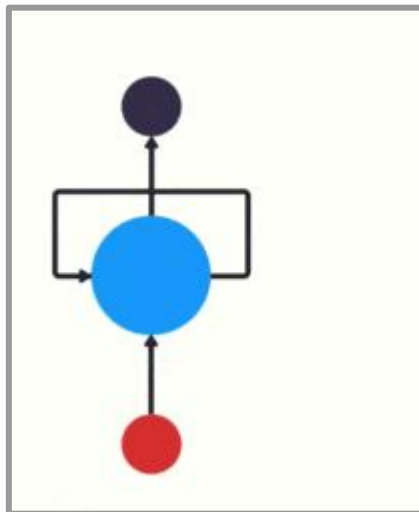
RNNs = Recurrent Neural Networks

Feed Forward



Fixed Input
Fixed Output
Independent Samples

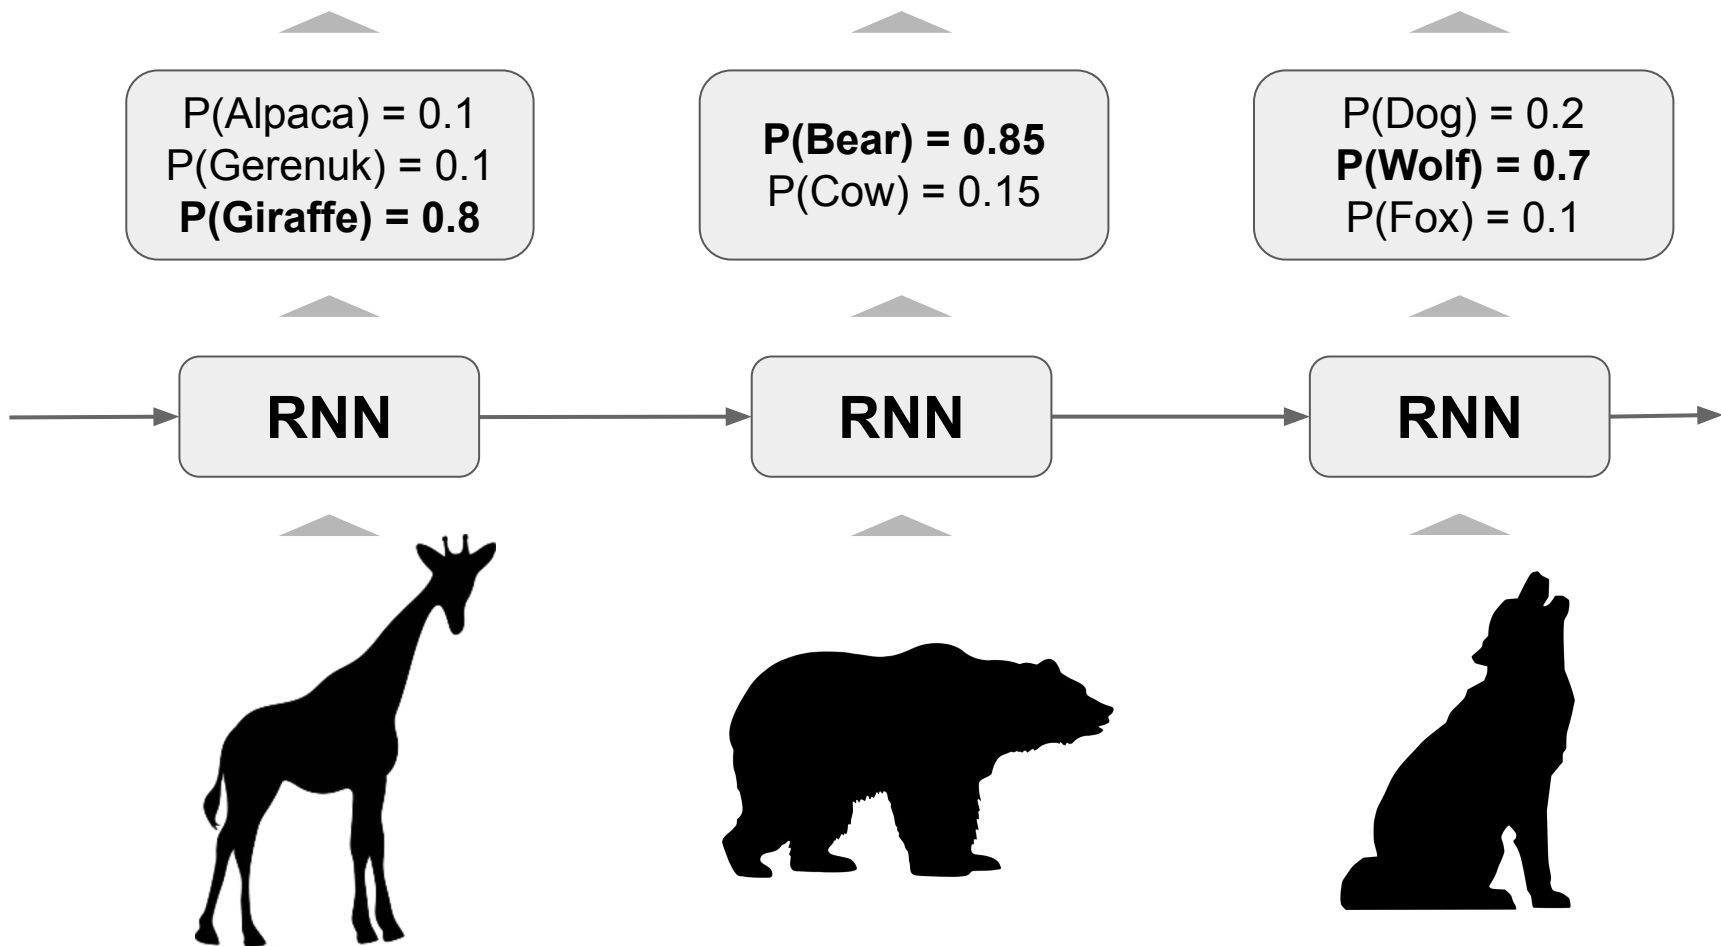
Recurrent



Unrolling RNN to Feed-Forward Networks



Dynamic Input
Dynamic Output
Dependent Samples



Strength of RNNs: sequence data predictions

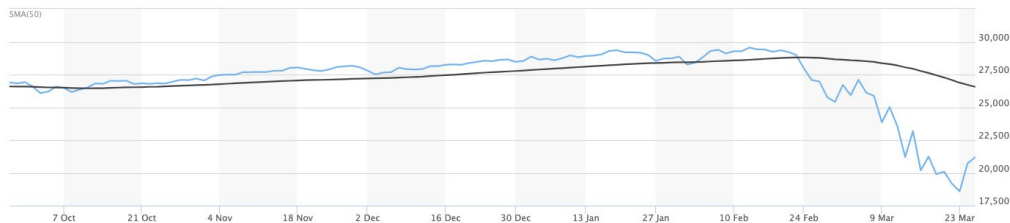
- Audio content in speech



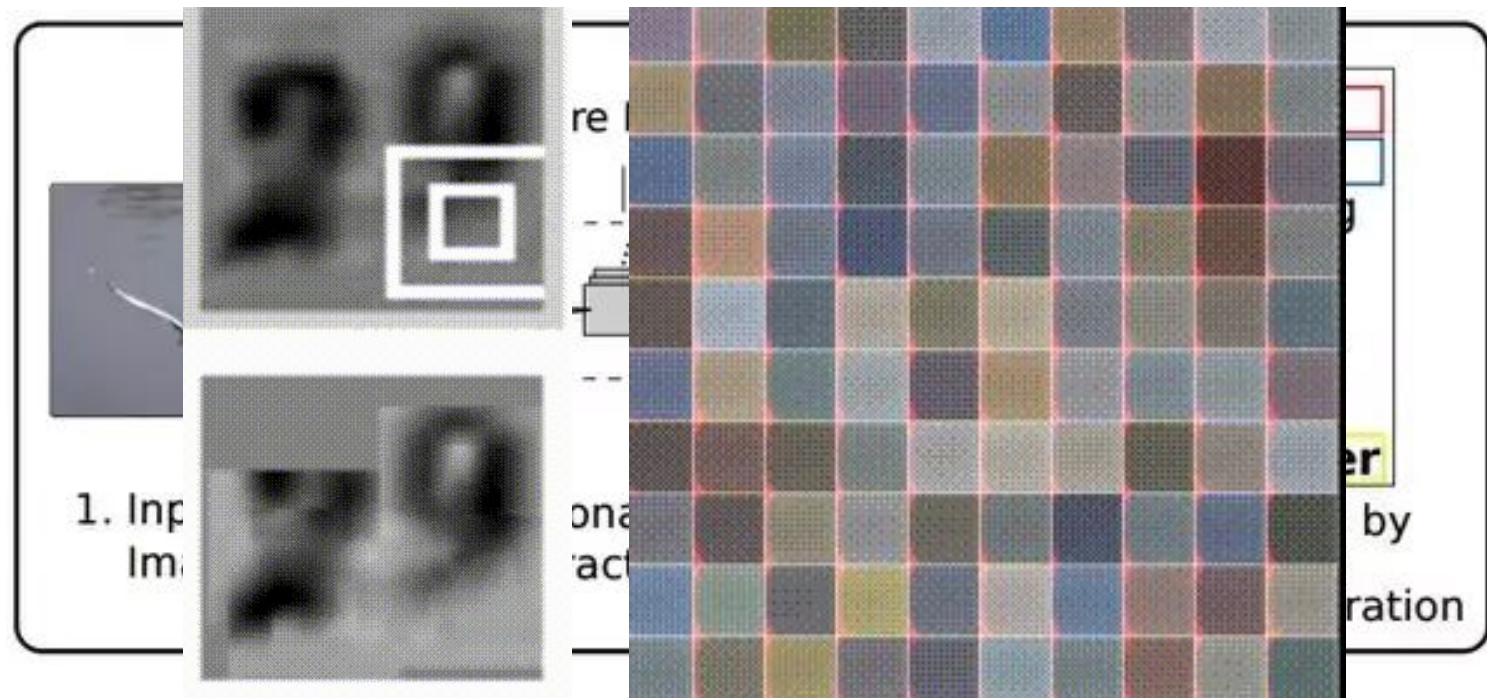
- Text content in articles

Text can also be sequence data

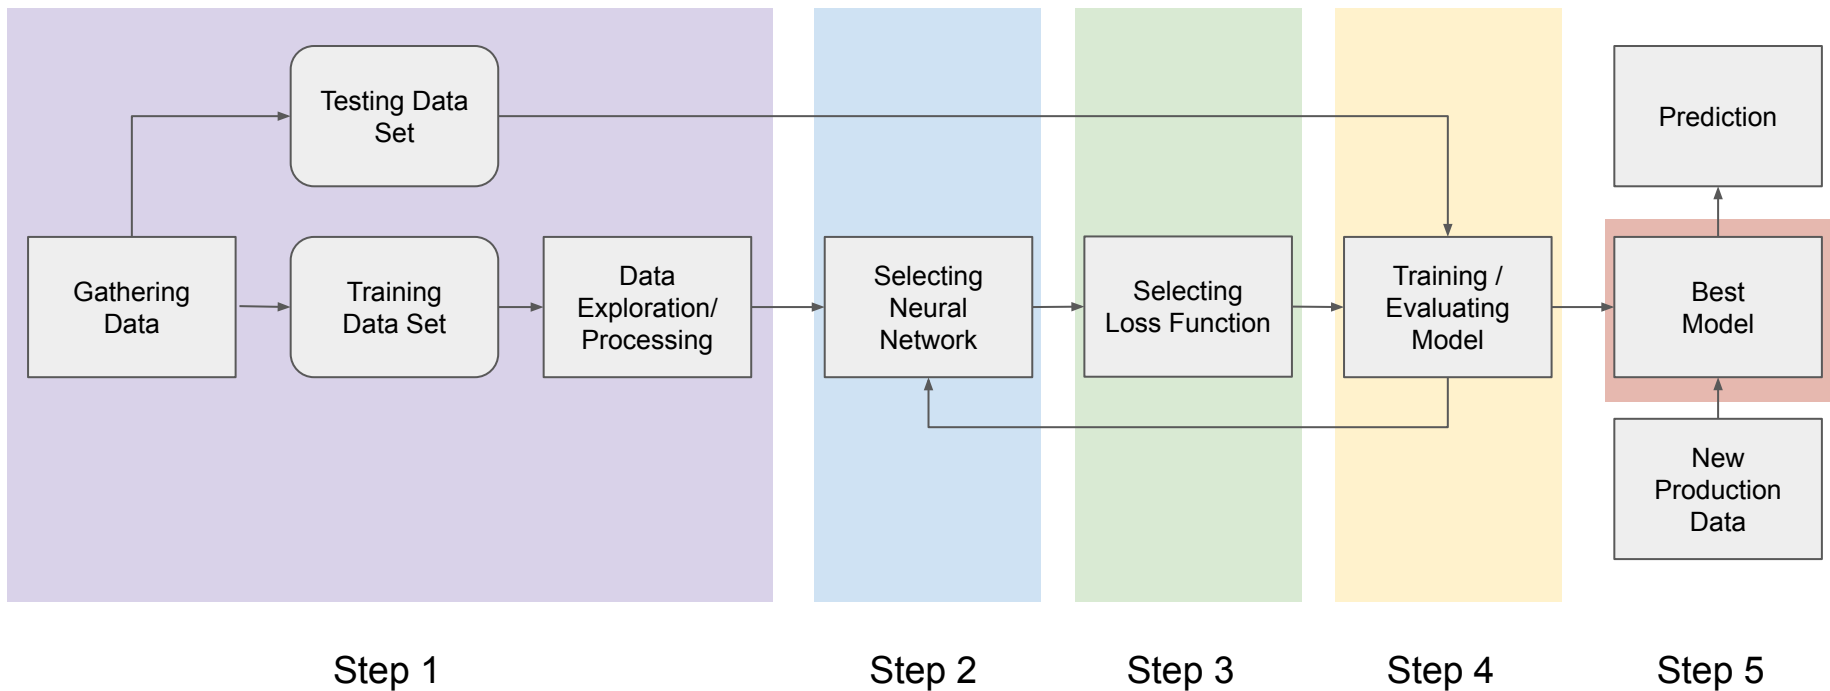
- Stock price prediction



Sequential processing in absence of sequences



Workflow for a deep learning project



Sequence Batching

[1 2 3 4 5 6 7 8 9 10 11 12]

Batch #1 → [1 2 3 4 5 6] Batch size = 2
Batch #2 → [7 8 9 10 11 12] Seq length = 6

OR

Batch #1 → [1 2 3] [4 5 6] Batch size = 2
Batch #2 → [7 8 9] [10 11 12] Seq length = 3

In our demo

[1 2 3 4 5 6 7 8 9 10 11 12]

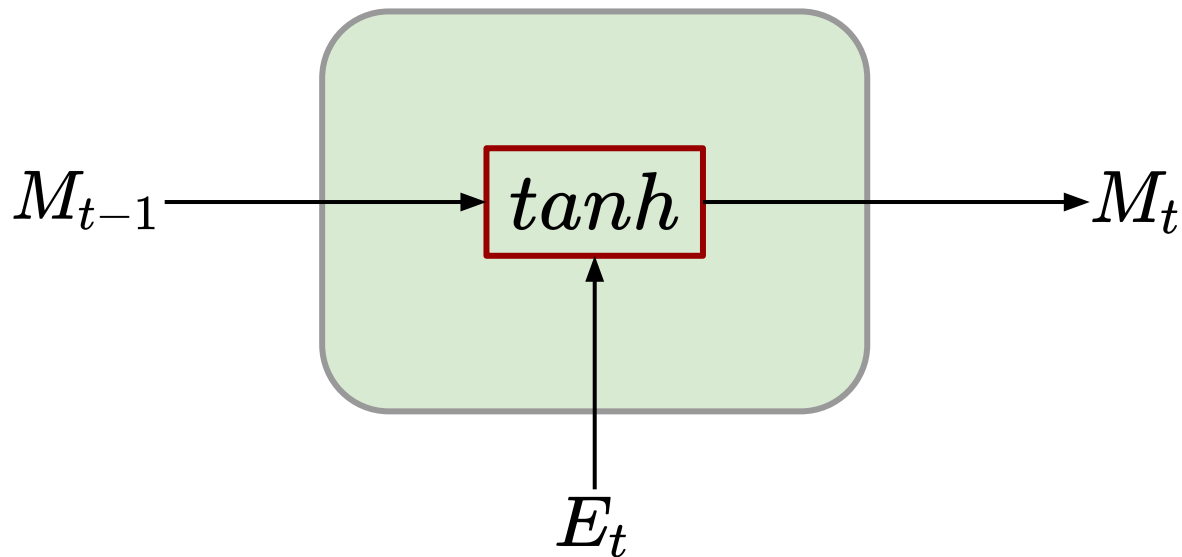
$$\begin{bmatrix} 1 & 2 & 3 & 4 & 5 & 6 & 7 & 8 \\ 2 & 3 & 4 & 5 & 6 & 7 & 8 & 9 \\ 3 & 4 & 5 & 6 & 7 & 8 & 9 & 10 \\ 4 & 5 & 6 & 7 & 8 & 9 & 10 & 11 \\ 5 & 6 & 7 & 8 & 9 & 10 & 11 & 12 \end{bmatrix}$$

Batch size = 5

Seq length = 8

RNN Inside

$$M_t = \tanh(W_{iM}E_t + b_{iM} + W_{MM}M_{t-1} + b_{MM})$$



Different Applications of RNNs

one to many

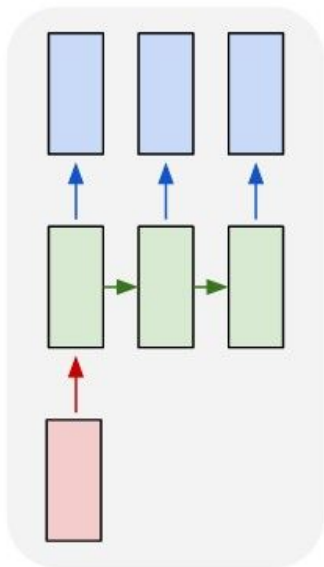
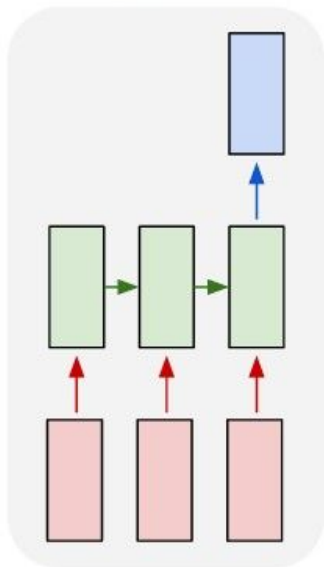


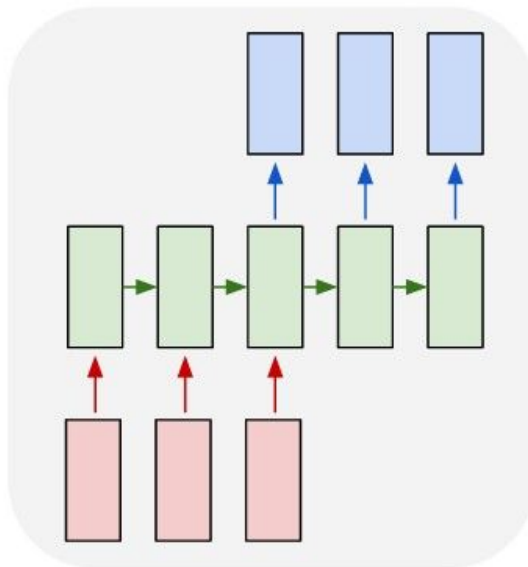
Image Captioning

many to one



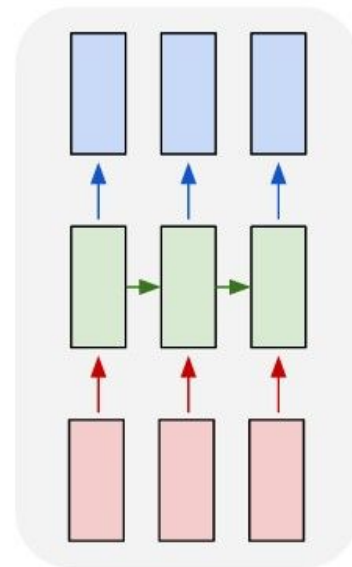
Sentiment Analysis

many to many



Machine Translation

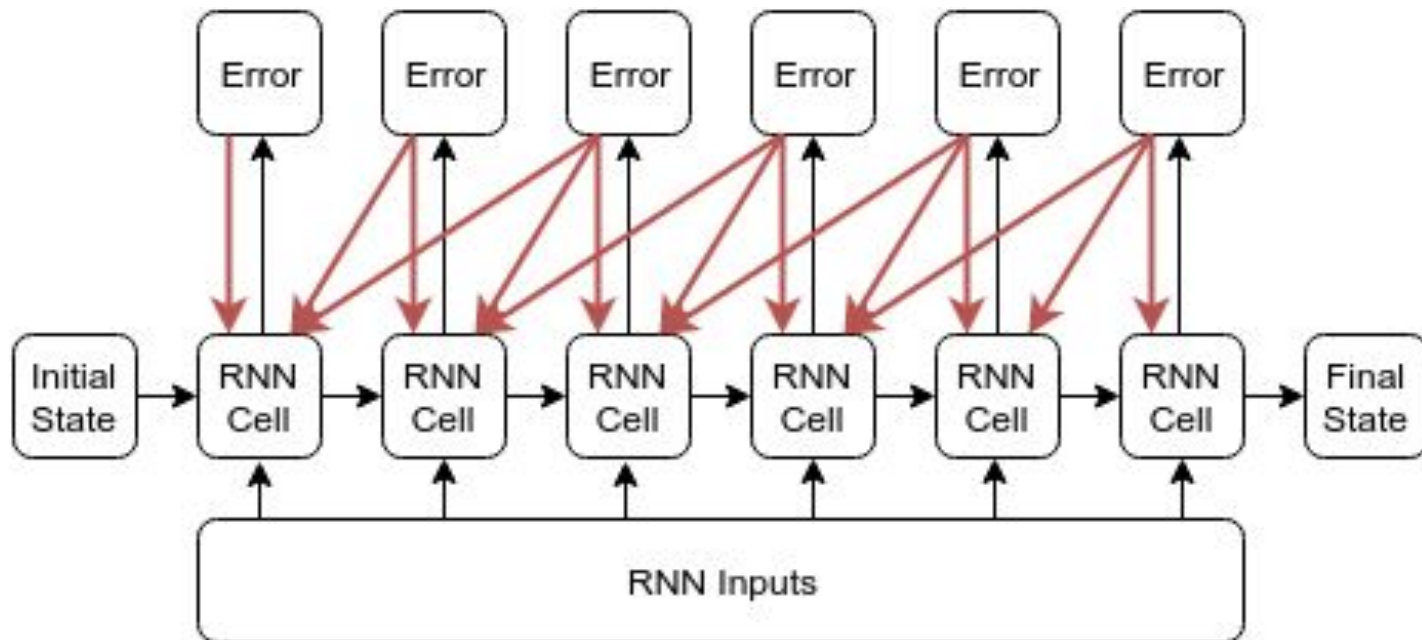
many to many



Video Classification

[Source](#)

Backpropagation Through Time (BPTT)



[Source](#)

RNN Computation in PyTorch

- Defining an RNN
 - Create an RNN layer
 - Add a fully-connected layer

- [torch.nn.RNN](#)

- input_dim
- hidden_dim
- output_dim
- num_layers
- Batch_first

- In forward function
 - Initialize hidden state
 - BPTT
 - Shape output

```
class RNN(nn.Module):
    def __init__(self, input_dim, hidden_dim, num_layers, output_dim):
        super(RNN, self).__init__()
        self.hidden_dim = hidden_dim

        # Number of hidden layers
        self.num_layers = num_layers

        # Building your RNN
        # batch_first=True causes input/output tensors to be of shape
        # (batch_size, seq_length, feature_dim)
        self.rnn = nn.RNN(input_dim, hidden_dim, num_layers, batch_first=True)

        # Readout layer
        self.fc = nn.Linear(hidden_dim, output_dim)

    def forward(self, x, hidden):
        # get RNN outputs
        out, hidden = self.rnn(x, hidden)

        # Index hidden state of last time step
        out = self.fc(out[:, -1, :])
        return out
```

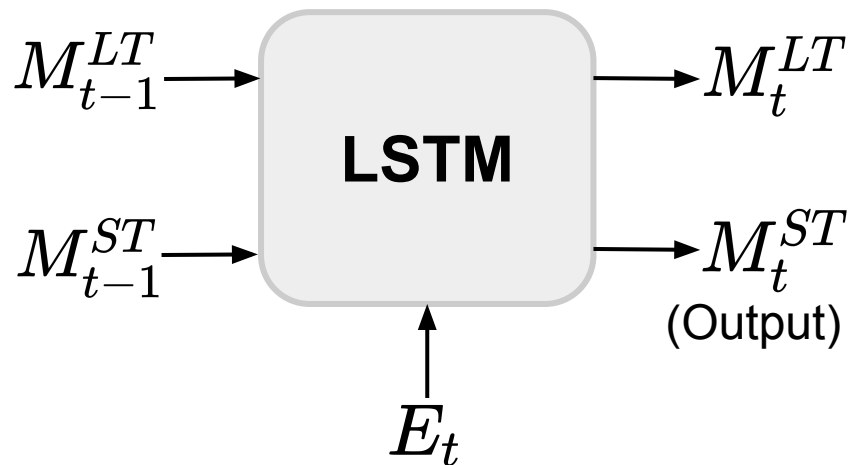
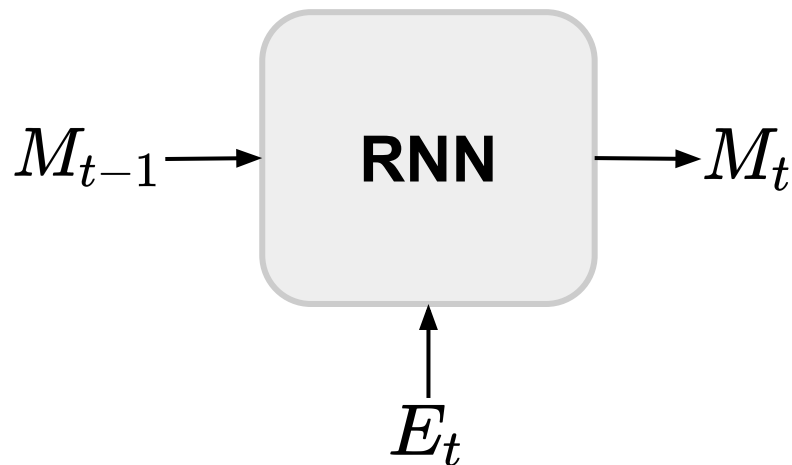
Colab Hands-on

bit.ly/LDL_04

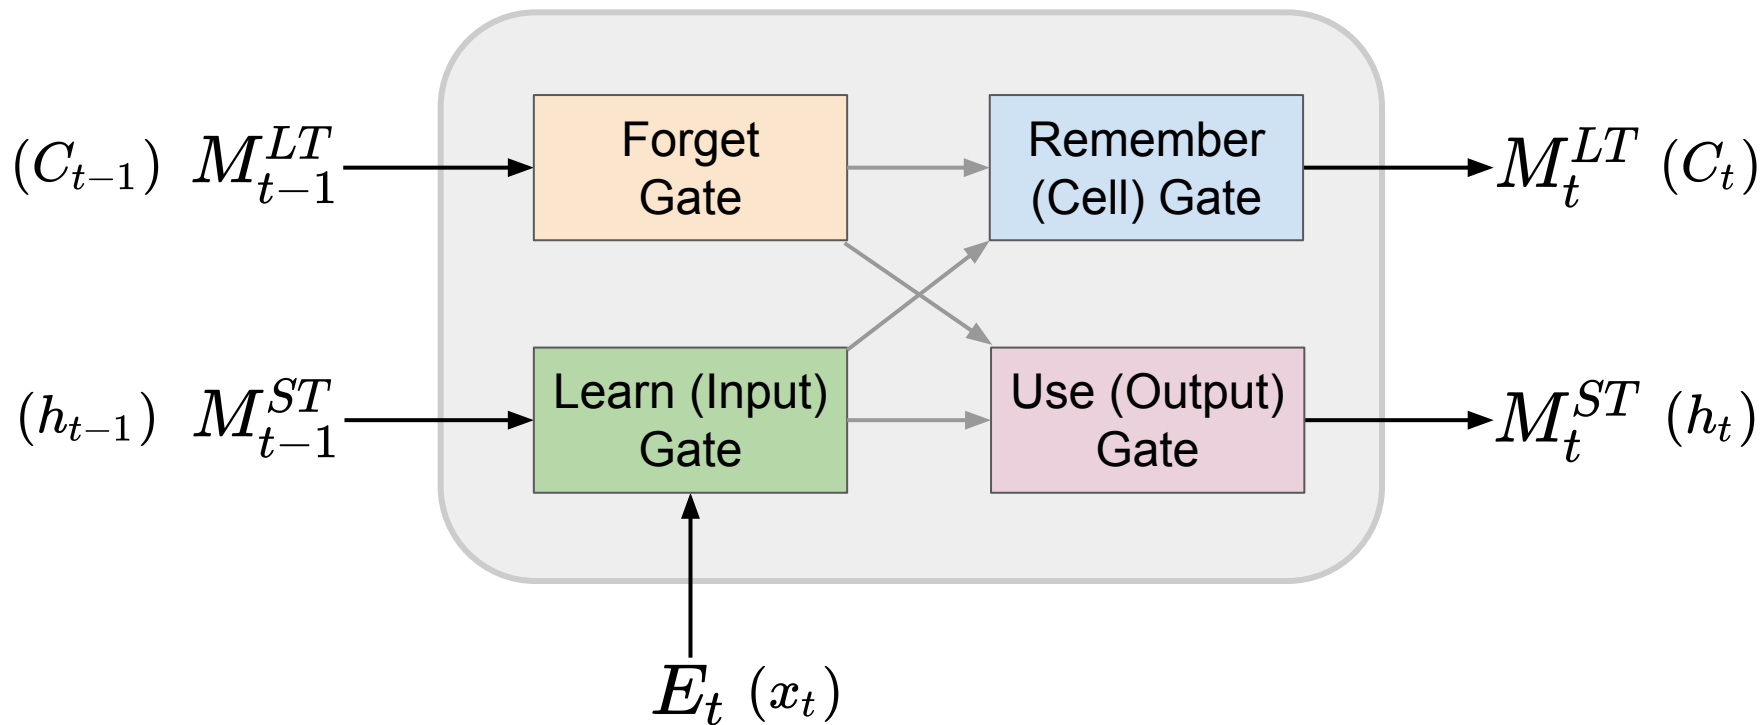
Problem in Vanilla RNN: short term of memory

What time is it?

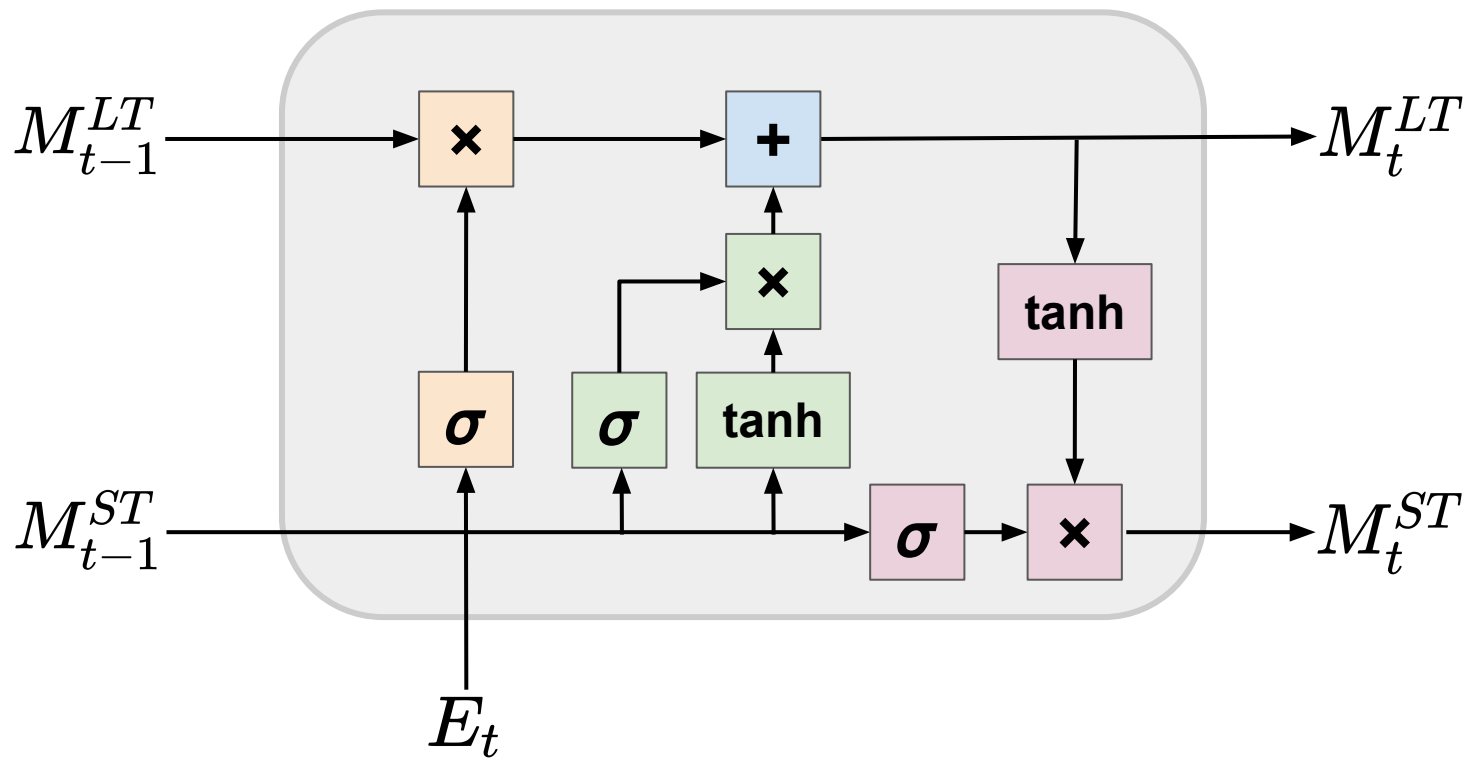
From RNN to Long Short-Term Memory (LSTM)



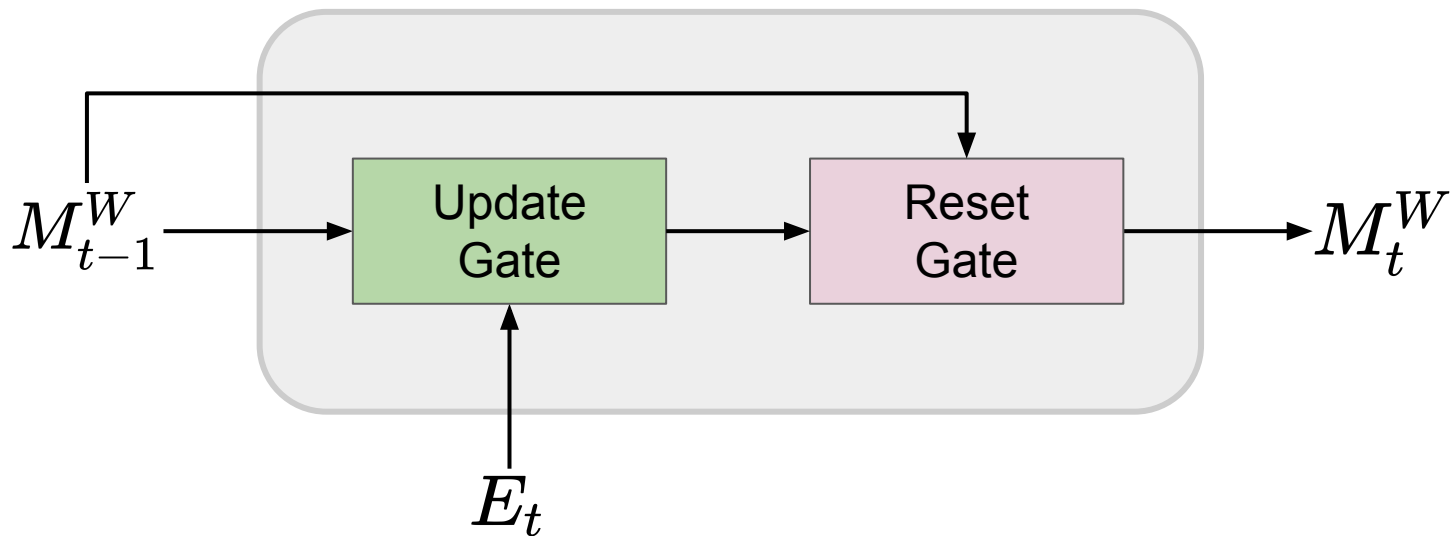
Basic Ideas of LSTM



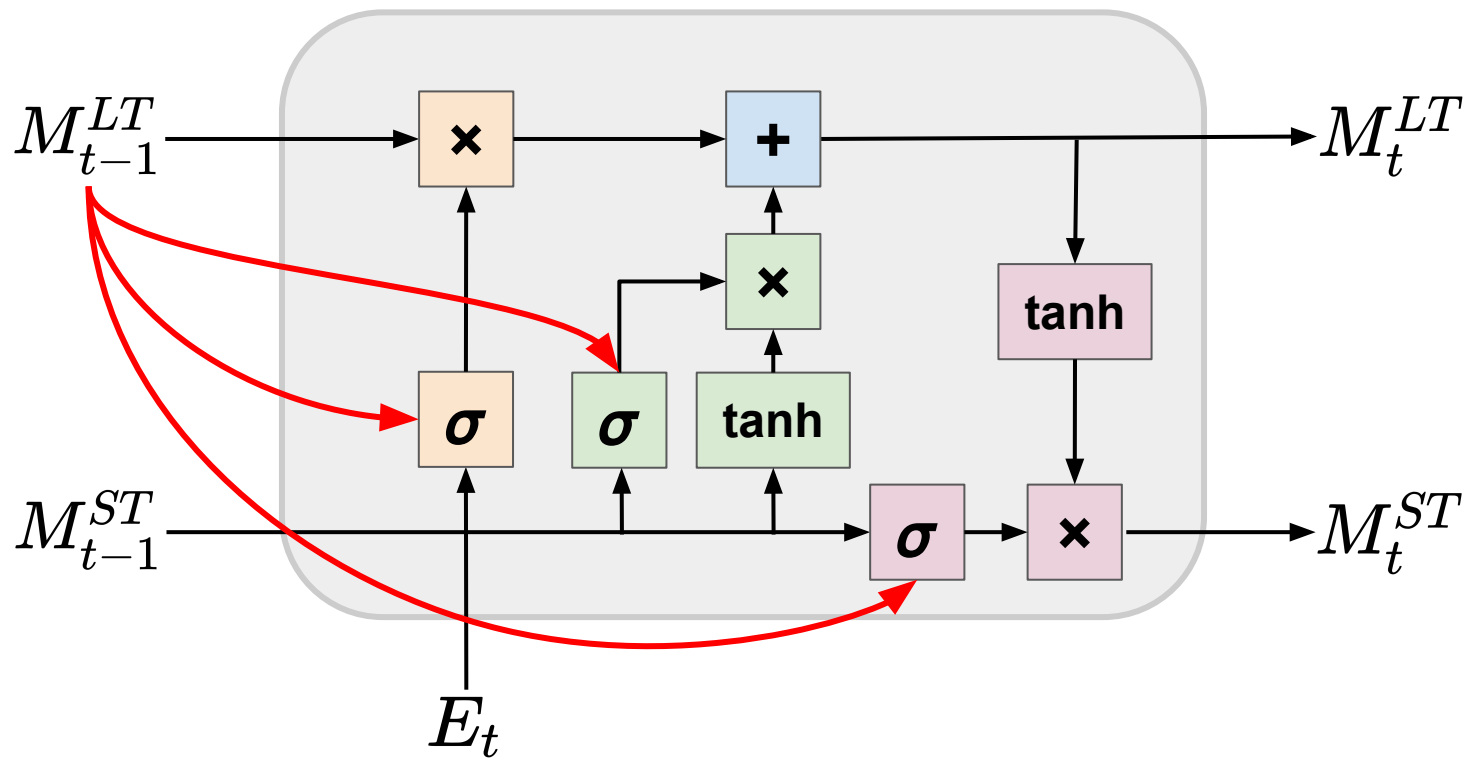
Architectures of LSTM



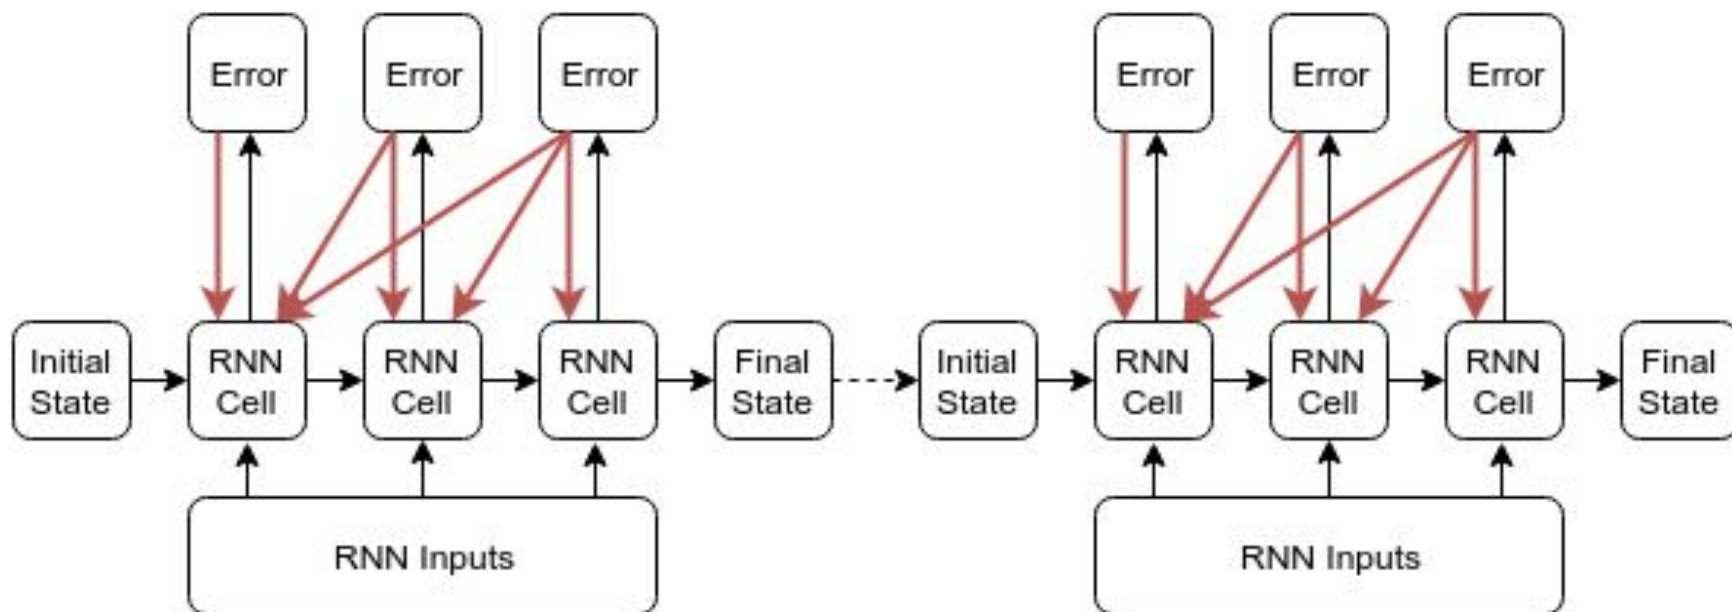
Other architectures: Gated Recurrent Unit (GRU)



Other architectures: Peephole LSTM

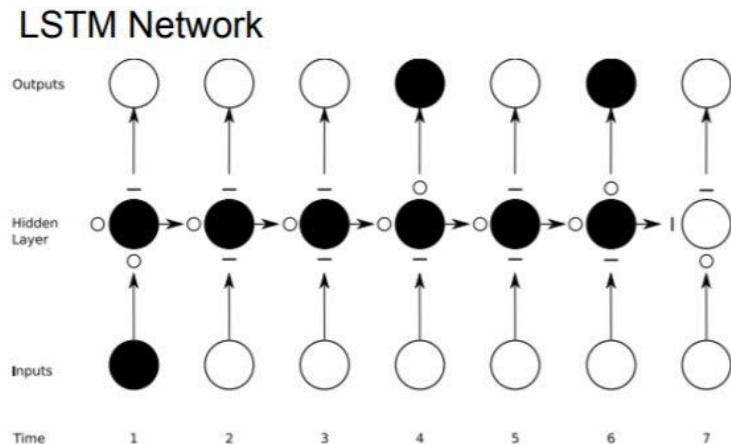
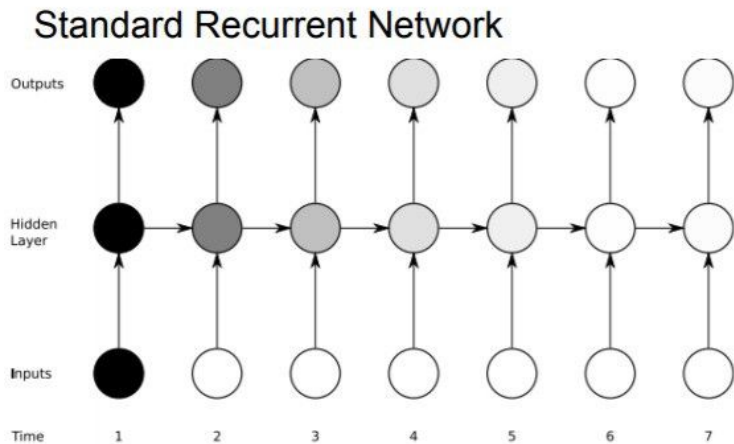


Truncated Backpropagation Through Time (TBPTT)



[Source](#)

LSTMs reduce vanishing gradient problem



Graves et al 2013

- The darker the shade, the greater the sensitivity
- The sensitivity decays exponentially over time as new inputs overwrite the activation of hidden unit and the network 'forgets' the first input

Colab Hands-on

bit.ly/LDL_04

More words on time-series predictions

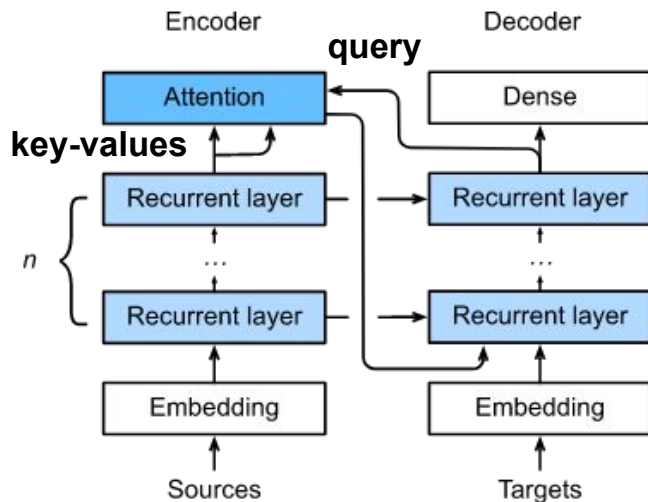
- Inspect the uncertainty in the series
 - Stochastic (truly random)
 - Epistemic (process too complex, looks like random)
- ARIMA-type model:
 - Implicit Gaussian assumptions
 - Parametric: to predict an individual time series (e.g. p , d & q)
 - Better for small dataset
 - Better results in forecasting short term
 - Python module statsmodels: `from statsmodels.tsa.arima_model import ARIMA`
- RNN-based model:
 - “Non”-parametric: possibly for unknown series
 - Better for large dataset
 - LSTM: better results in forecasting long term

Attention is all you need!

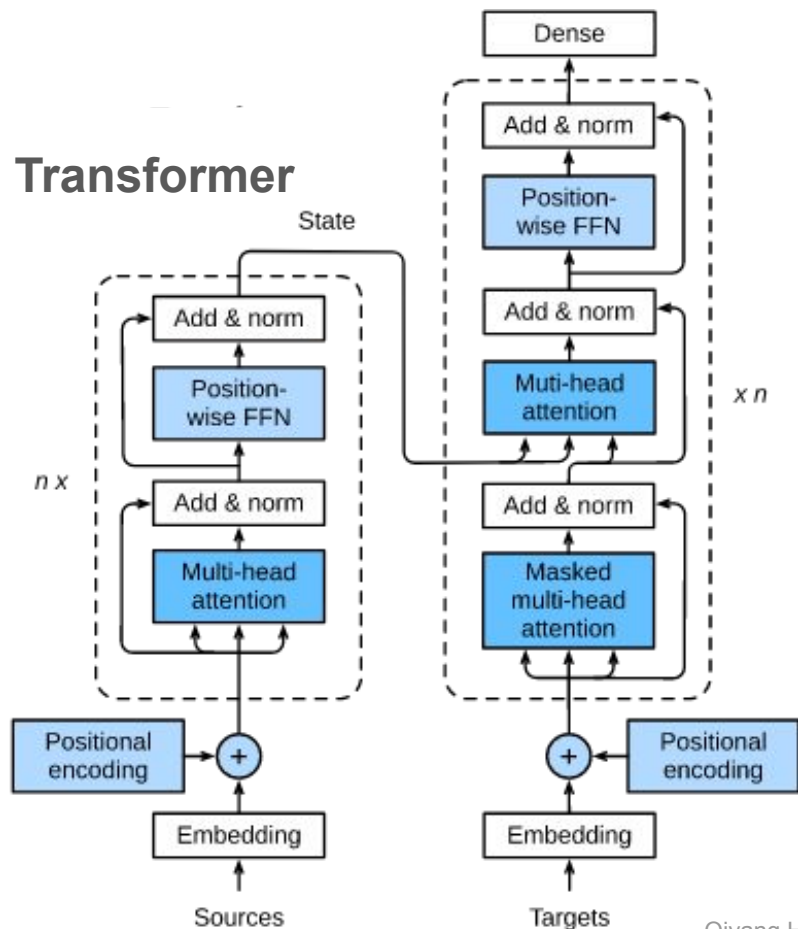
- **Seq2Seq**

- Introduce weights for encoding outputs

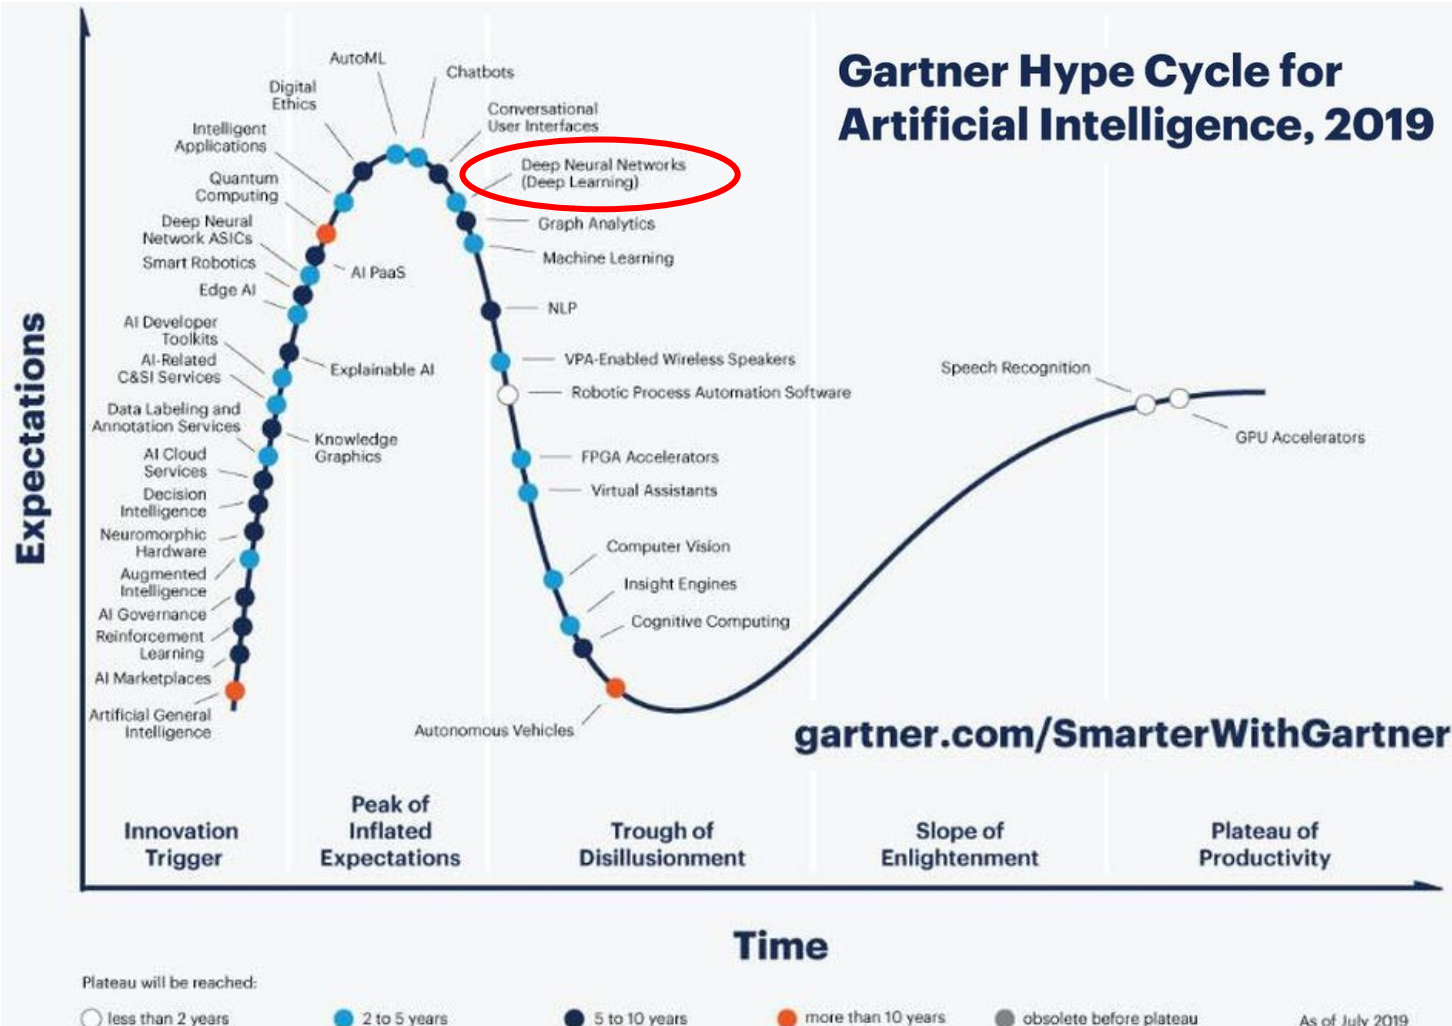
$$r = \sum_{i=1}^n \alpha(k_i, q) \cdot v_i$$



- **Transformer**



Gartner Hype Cycle for Artificial Intelligence, 2019



We want to hear from you!

bit.ly/2X2phyS

- Contact me for questions or discussions
 - hugy@idre.ucla.edu
 - Office: Math Sci #3330
 - Phone: 310-825-2011